

# Introduction To SeaScape

2020 R3 Release



# Guidelines For Training

- This event will be for 2 hours. There will be 15 min Q&A at the end of this session for questions.
- To ask questions during training, use the Q&A window and select option "All panelists" to ask question
- For audio, join using Audio Broadcast option
- Slides and recording links will be sent in 1-2 days after training
- Users successfully completing the training will get a digital certificate.
  - Need minimum 70 % attendance for this ( 10/14 training sessions )

# Agenda



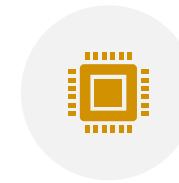
Introduction To  
Big Data



SeaScape  
Platform



User Interface



SeaScape Help  
System & API's



Mapreduce  
Programming



Using SeaScape  
Platform

# **Introduction To Big Data**



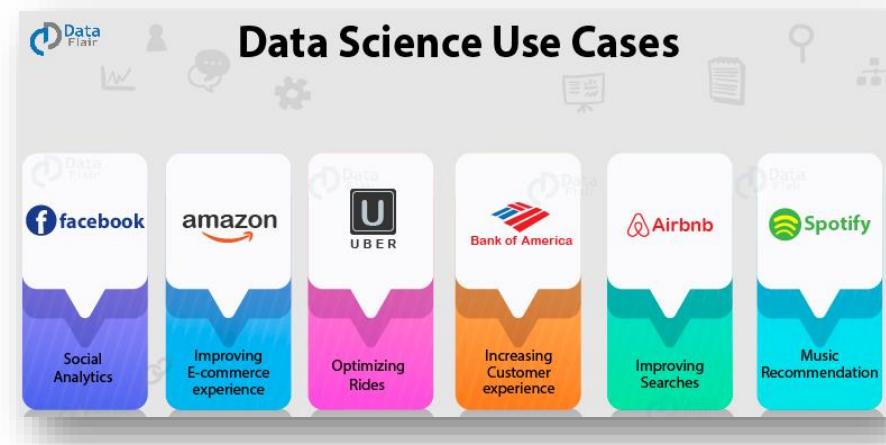
# Introduction to Big Data Analytics



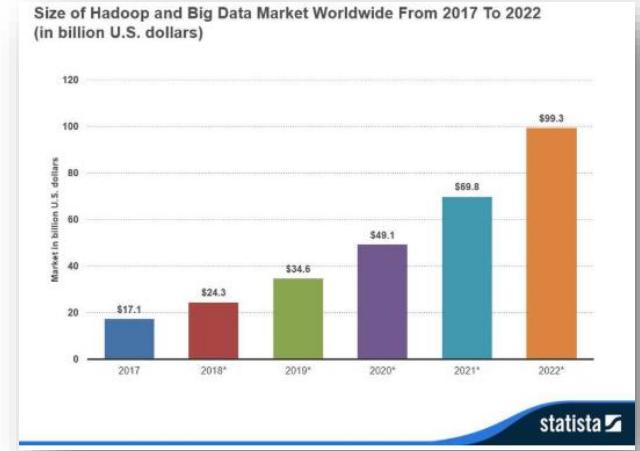
How big data analytics help ?  
Image courtesy : <https://sas.com>



Challenges in Big Data  
Image courtesy : <https://edureka.com>



Top Applications of Big Data  
Image courtesy : <https://data-flair.training>



Growing market of big data science – 100 B dollar by 2022  
Image courtesy : [forbes.com](https://forbes.com)

# Big Data Techniques

Google search says :

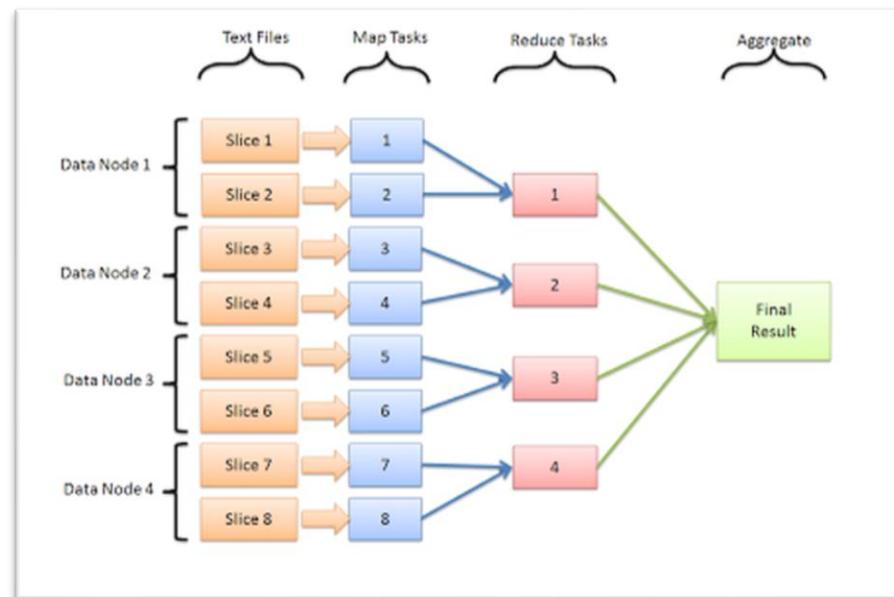
**MapReduce** is the heart of **Apache® Hadoop®**. It is this programming paradigm that allows for massive scalability across hundreds or thousands of servers in a **Hadoop** cluster

Ansys says :

**Scheduling algorithms similar to MapReduce** are the heart of **SeaScape®**. It is this programming paradigm that allows for massive scalability across hundreds or thousands of servers in **customer** cluster

# Big Data Techniques

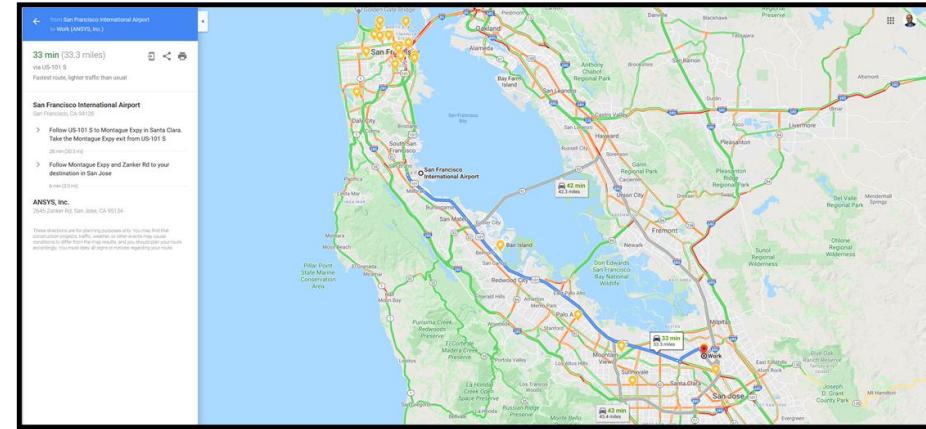
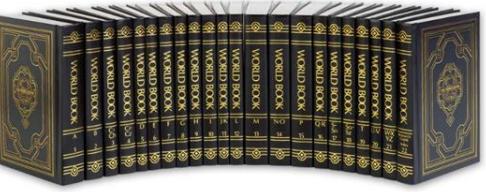
- **MapReduce** is a simplified model for processing large data sets with a parallel, distributed algorithm on a cluster in an automated, easy to write, fashion
- All big-data engines use MapReduce or its derived/augmented techniques
- Some popular big data engines below:



# SeaScape Platform



# Foundational Innovation was Necessary



## “Pre-Google™”

The data was there, but you had to work hard to make it useful.

## Today

Data is instantly available and useful; freeing you to create more

# Traditional EDA Big Data Exists in Silos



## Problem with Silos

Design margins too large

Coverage is poor

Results are not actionable

Multi-Domain Chip-Level Analysis is not well supported by existing EDA tools and databases

Search  
Maps  
Big data  
Hadoop

+

1 million vectors  
1 billion logical instances  
10 billion transistors  
100 billion geometries  
1 trillion model elements

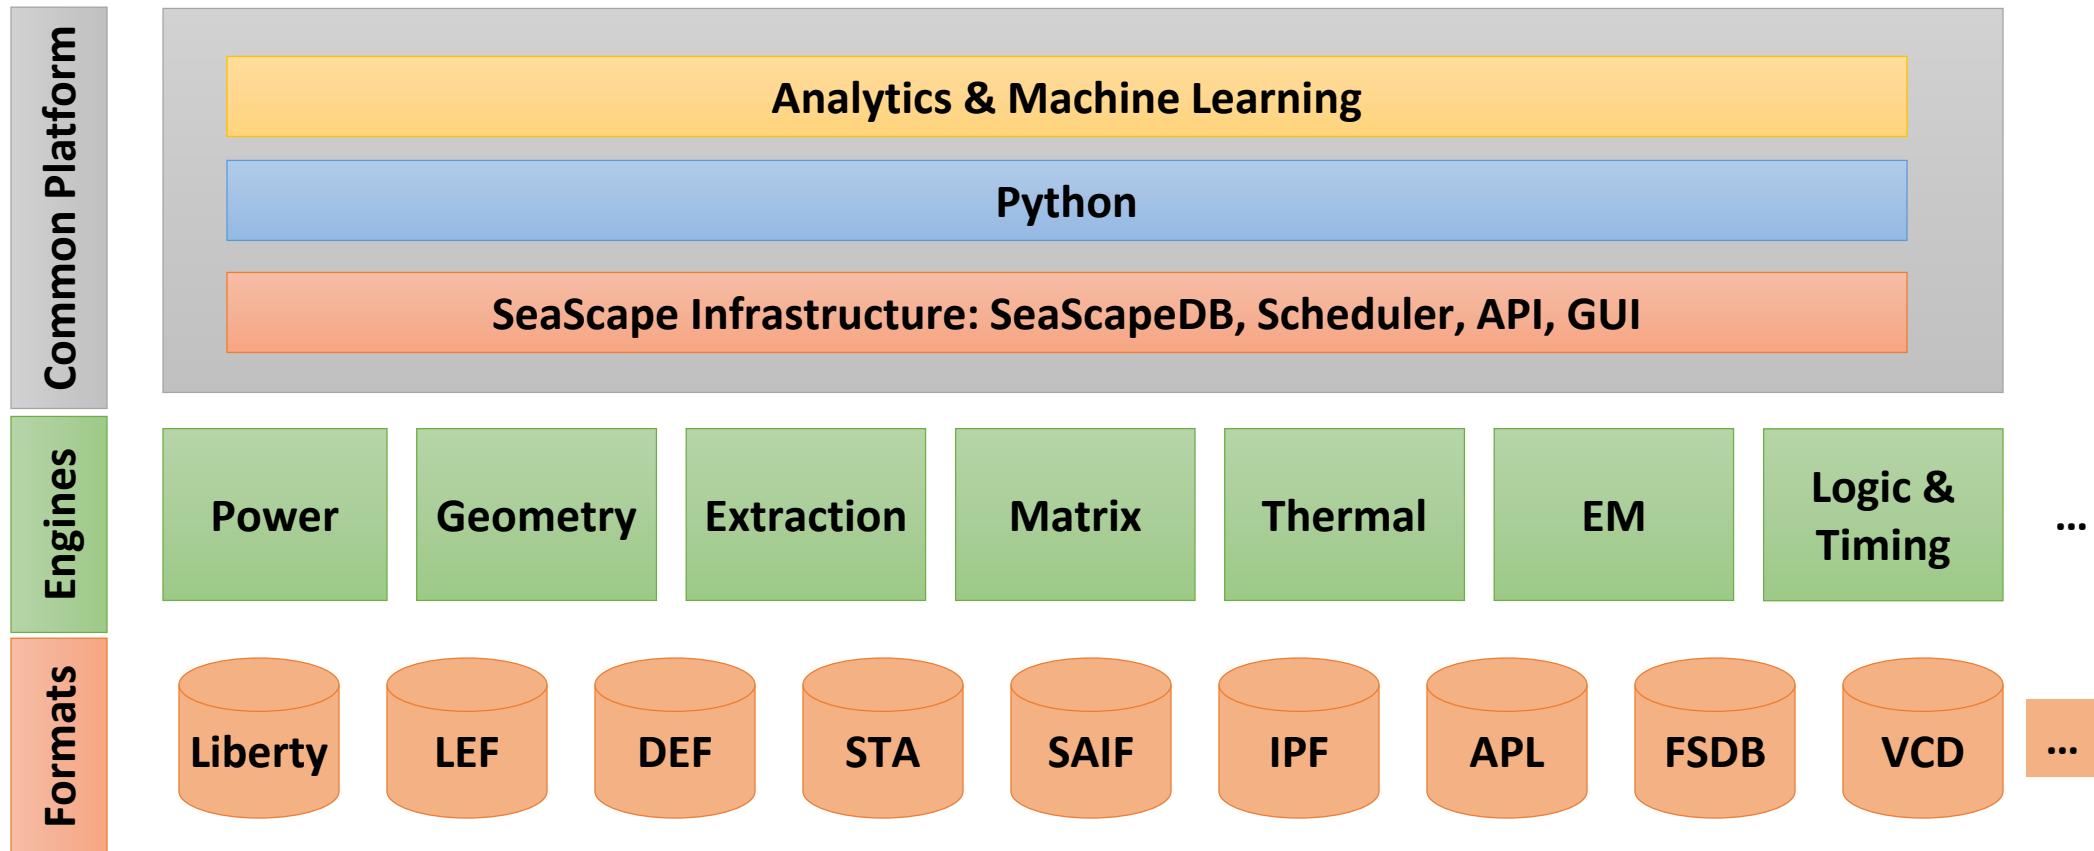
= ?

Big data for chip design?

# The SeaScape Platform

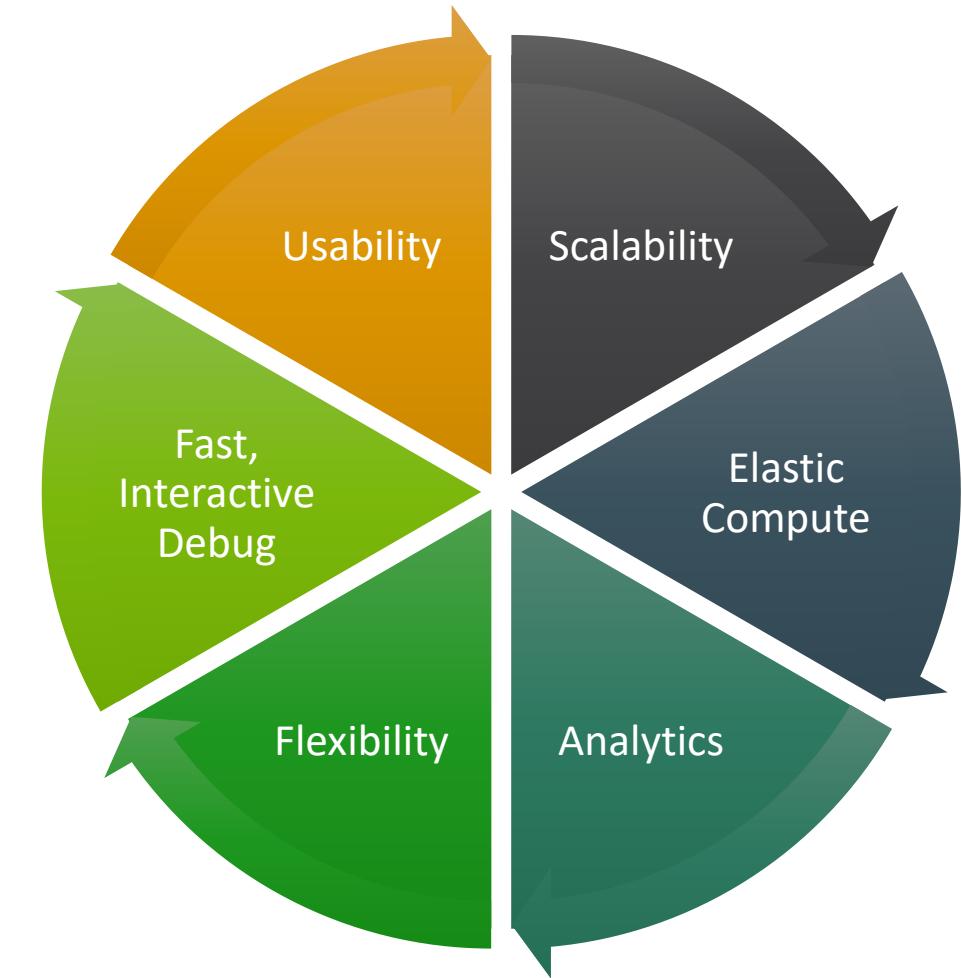
*EDA-specific platform for distributed data and compute*

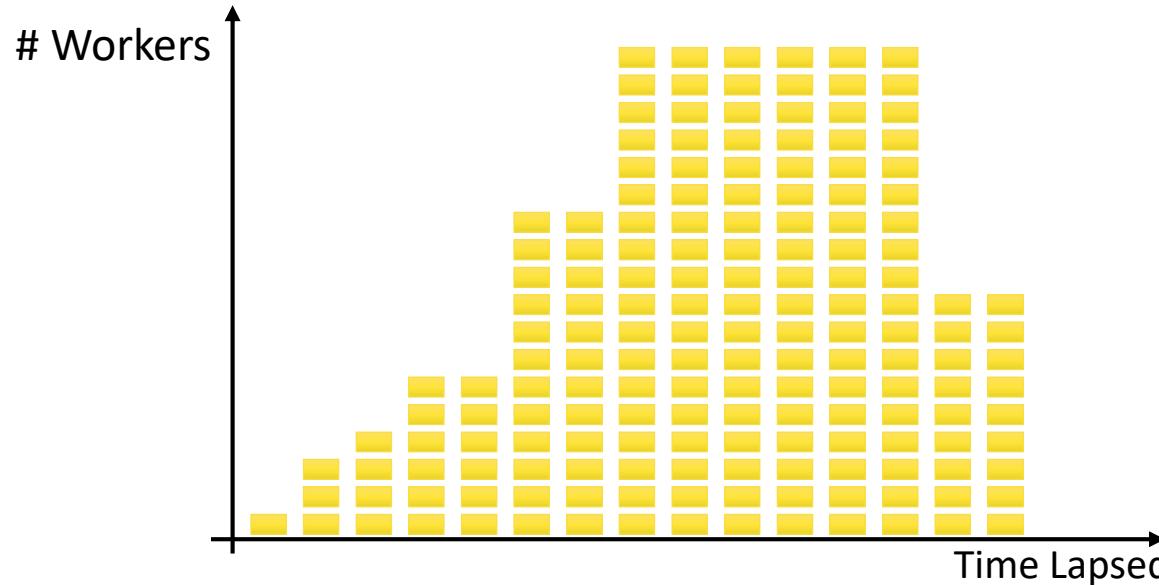
1. SeaScape is the industry's first big data platform for Actionable analytics, using Big data and machine-learning apps
2. All products of ANSYS will work with the SeaScape platform.



# Advantages of using SeaScape Platform

- **Scalability**: can handle massive full-chip designs with billions of instances
- **Elastic Compute**: efficient use of hardware resources
- **Powerful Multi-Variable Analytics**: User can apply MapReduce to query data and compute custom results, generate flexible heatmaps
- **Flexibility**: one run can open multiple dbs, each with multiple views
- **Fast Interactive Debug**: load huge design layout and results in multiple GUI in seconds
- **Usability**: MapReduce Wizard, Console, Online Help





### Launch runs immediately

- No waiting for full machines
- Instantly view results on single cores

### Auto checkpoint

- Reuse data from terminated runs
- Build multiple flows & reuse data

### Resiliency

- Auto-recover from bad cores
- Or network issues

### Parallelization

- All stages are scalable across cores
- Matrix solve > 1,000 cores

## SeaScape Big Data Platform

Actionable Analytics

Machine Learning / AI

Elastic Compute

Visualization & Debug

# Launching Workers

- **Definition of worker :**

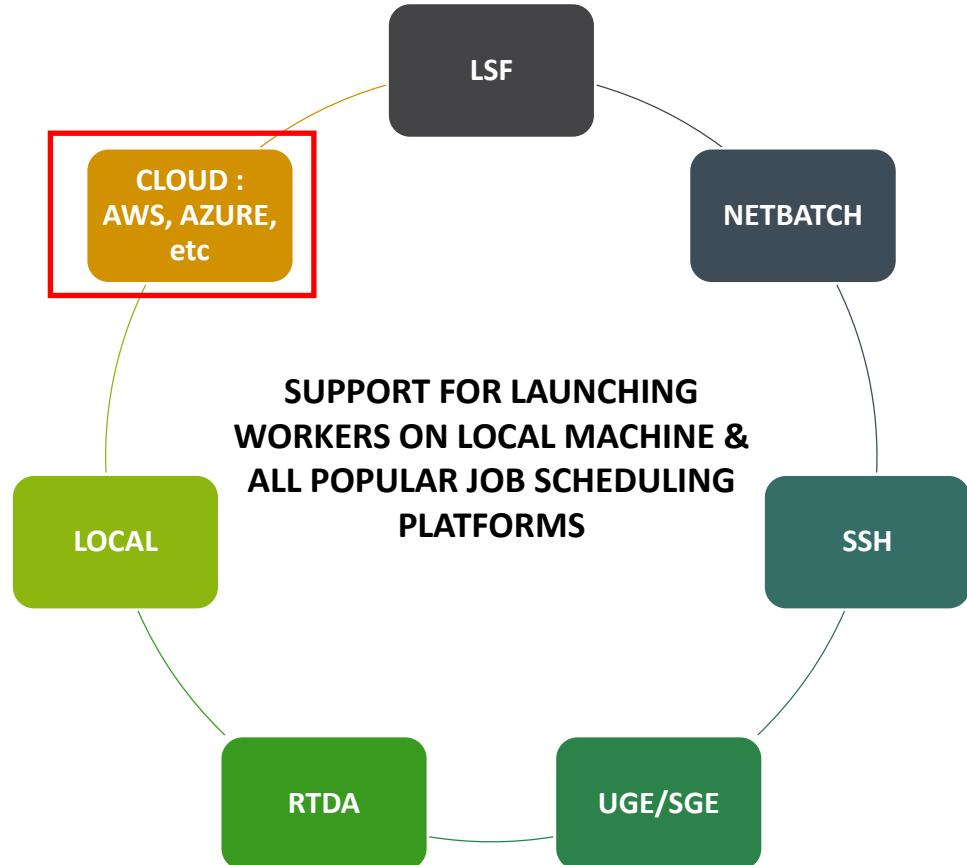
- SC is a distributed platform, runs on a set of CPU cores
- CPU cores used by RH-SC are called workers
- SC can launch hundreds of thousands of workers to process jobs efficiently
- 1 worker == 1 CPU core with certain amount of memory

- **Where are the workers located ?**

- Workers can be on the same machine or different machines in a network

- **How are workers launched ?**

- SC supports all popular job scheduling platforms like LSF , UGE , RTDA etc
- SC can also run on cloud based services like AWS, Azure etc



# More On Workers

- **Launcher commands syntax :**

```
create_grid_launcher(name, submit_command, num_workers_per_launch=None, initial_exec_function=None,  
max_num_threads_per_worker=None, memory_gb_per_worker=None)
```

- **name** : Launcher identifier name (type=str, required=True)
- **submit\_command** : grid submit launch command (type=str, required=True)
- **num\_workers\_per\_launch** : specify number of workers per launch (auto detect if not specified) (type=int, default\_value=None)
- **initial\_exec\_function** : function pointer (type=object, default\_value=None)
- **max\_num\_threads\_per\_worker** : maximum number of threads per worker (type=object, default\_value=None, constraint="int > 0 or string 'auto'")
- **memory\_gb\_per\_worker** : requested memory in GB (type=int, default\_value=None, constraint="int > 0")

- **Example of UGE launcher command :**

- `ll = create_grid_launcher('uuu','qsub -V -b y -q sjo_all_hosts -l mfree=10G -l h='sjoq*' -l platform='linux64e7' -o uge.log'. num_workers_per_launch=1, max_num_threads_per_worker=2)`

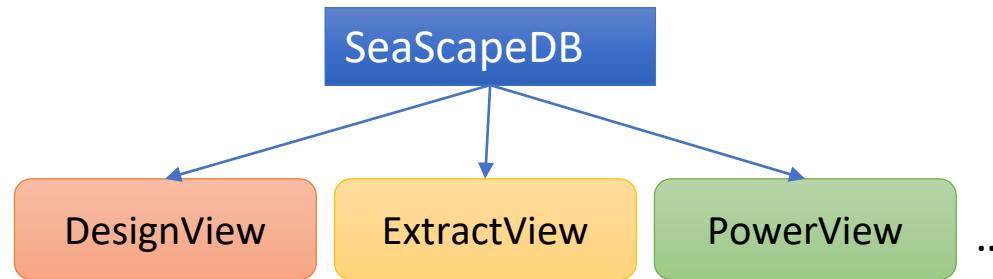
- **Launch workers:**

- `ll.launch(64) # launch 64 workers`

- **Making The Launcher Automatic (optional):**

- use 'register\_default\_launcher' function for SeaScape to automatically launch as many workers as needed [as run progresses]:
  - `register_default_launcher(launcher=ll , min_num_workers=5, max_num_workers=50)`

# The SeaScape Database



- SeaScapeDB holds all data used within SeaScape
- Data from external formats like DEF or FSDB
- Results of processing within the tool such as instance power

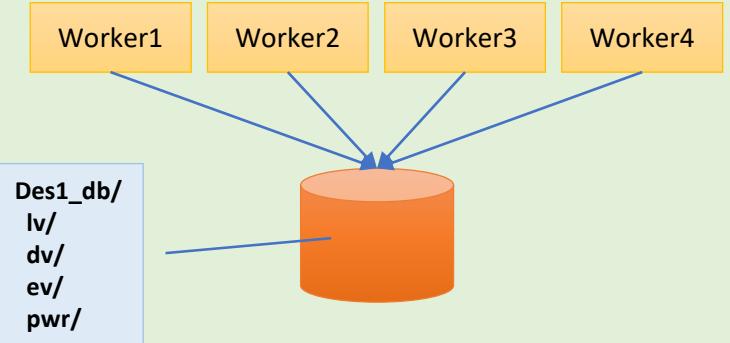
- One db contains **multiple views**
- Each view holds related data
- All data is partitioned into chunks for efficient processing

- Data is initially stored in each worker's memory, then later written to disk in the background

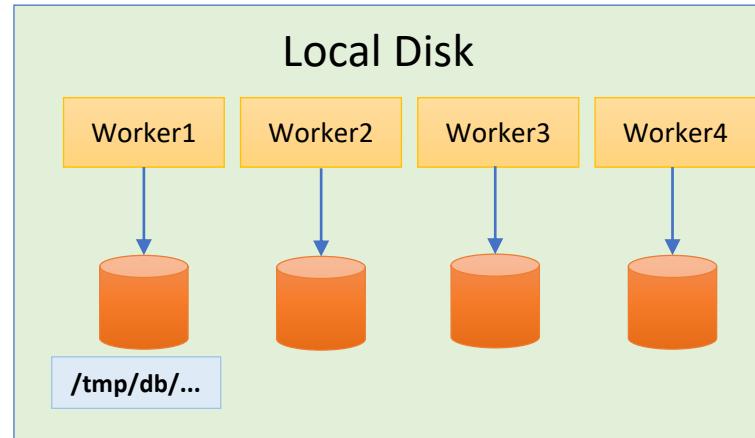
- Can read and/or write to multiple dbs in one tool session

## DB Storage Options

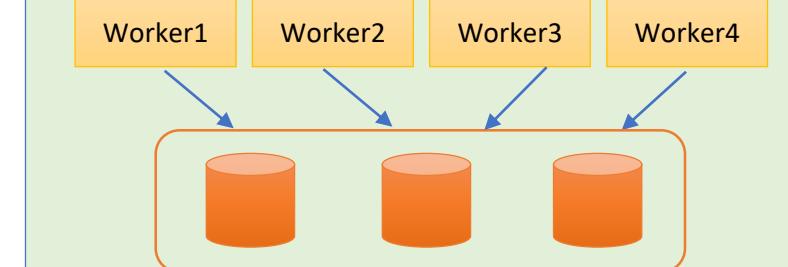
Central Disk



Local Disk

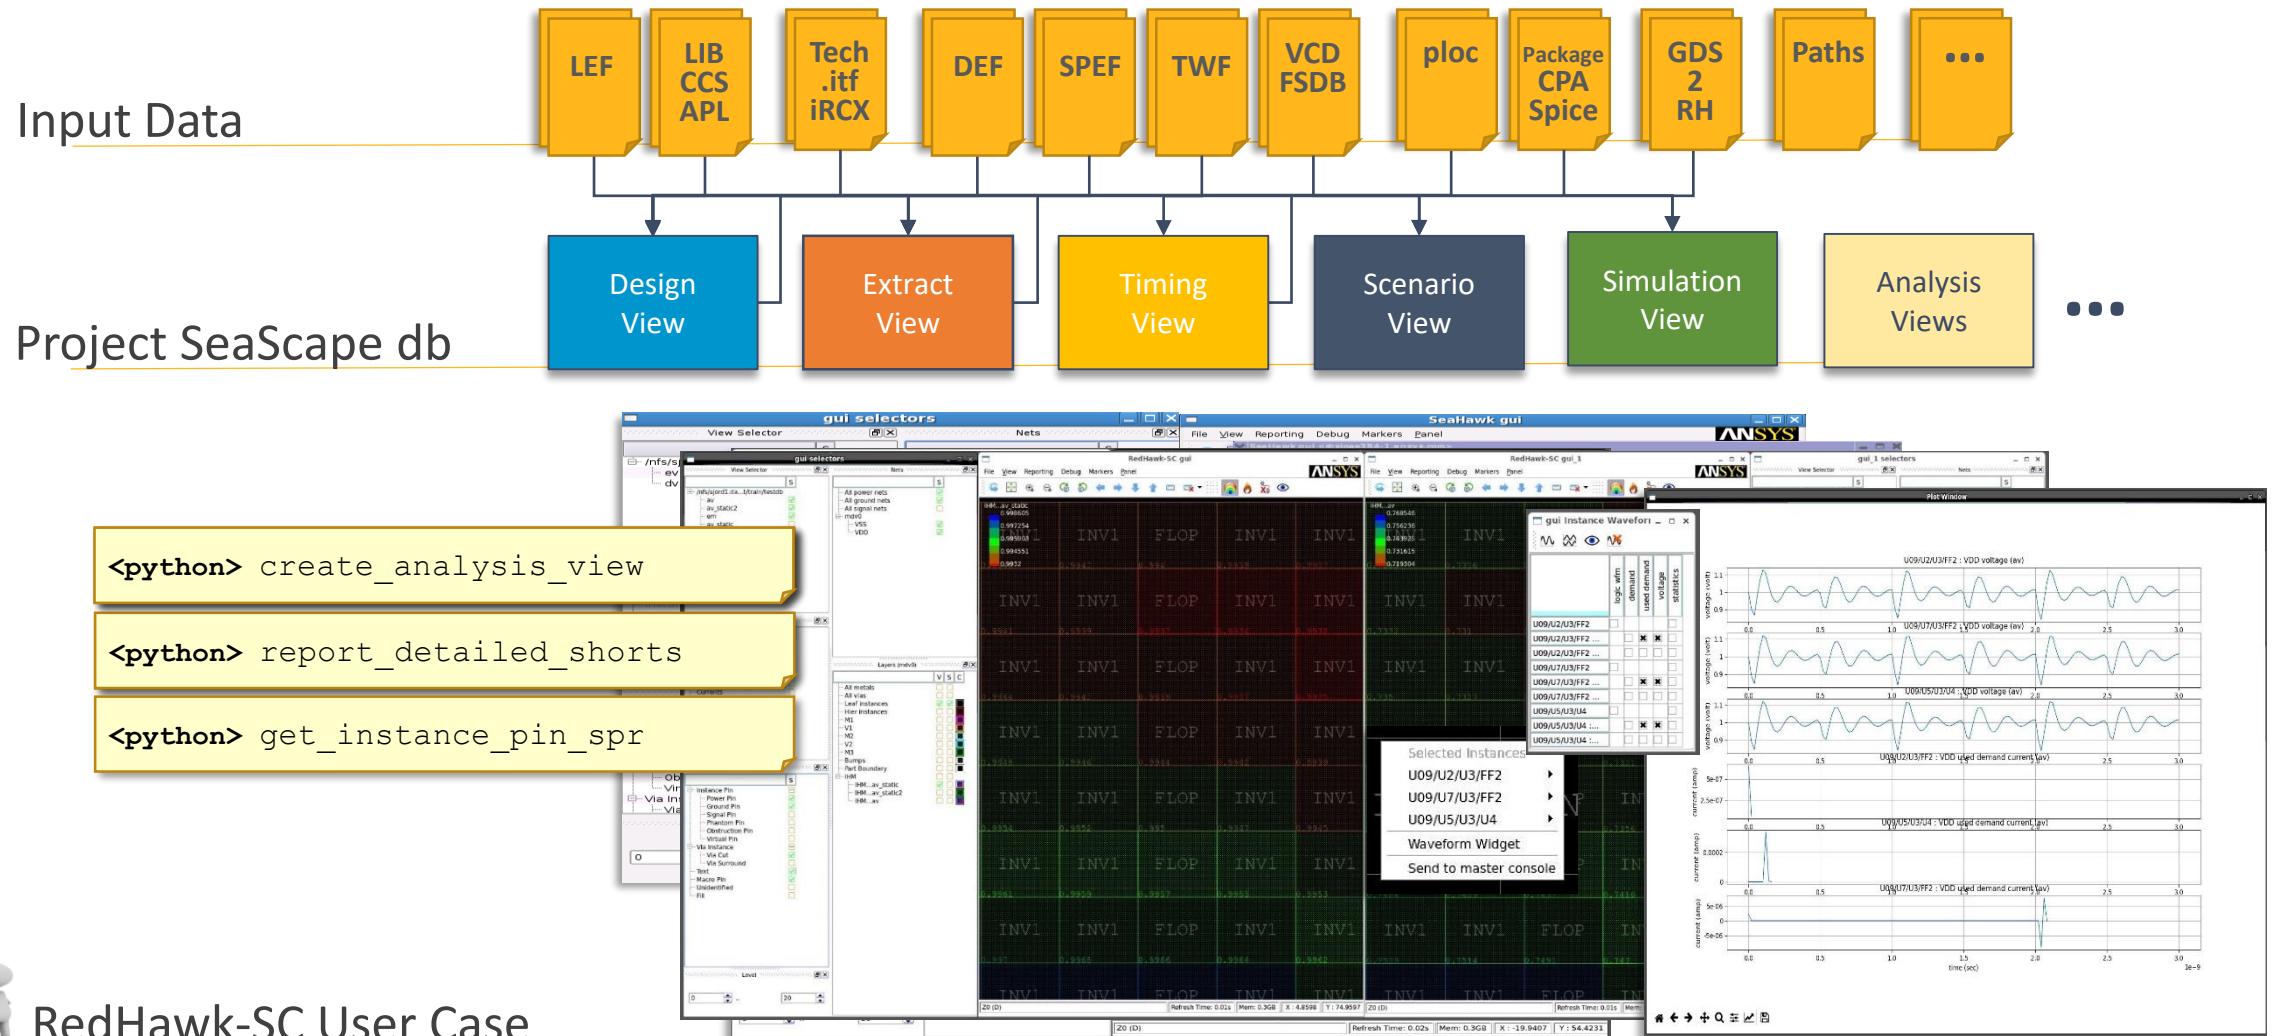


Multiple Central Disk Locations



# SeaScape

EDA-Project Data into SC “Views”



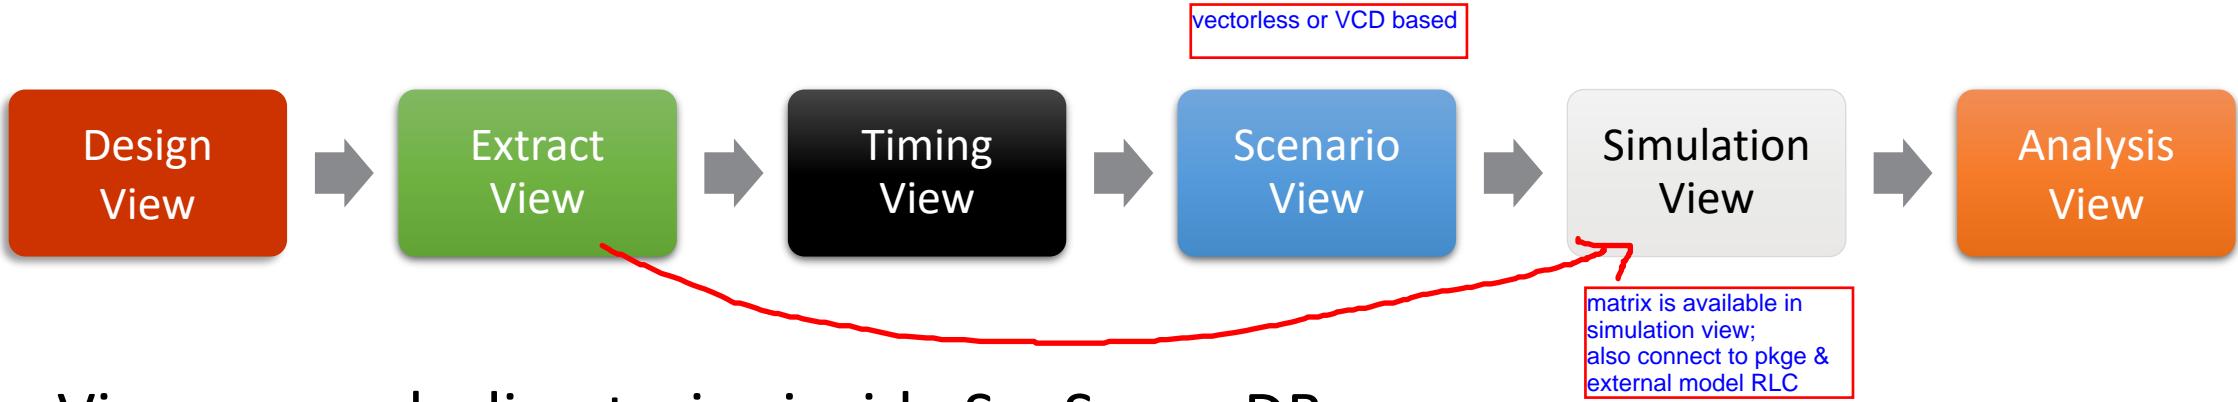
RedHawk-SC User Case

# SeaScapeDB – A View Oriented Data System

## SeaScape DB Overview

- A SeaScapeDB is simply a linux directory containing views
- The db is created through “*open\_db*” command
- Directory name for the DB is set in the “*open\_db*” command
- If db already exists, the db will be opened
- DB contains several views, each view is a sub-directory in the DB
- Views are created through “*create\_<type>\_view*” command
- Directory name for the view can be specified while creating the view
- Multiple SeaScapeDBs can be opened simultaneously
- A DB can be accessed simultaneously from multiple RH-SC sessions

# Views in SeaScape DB

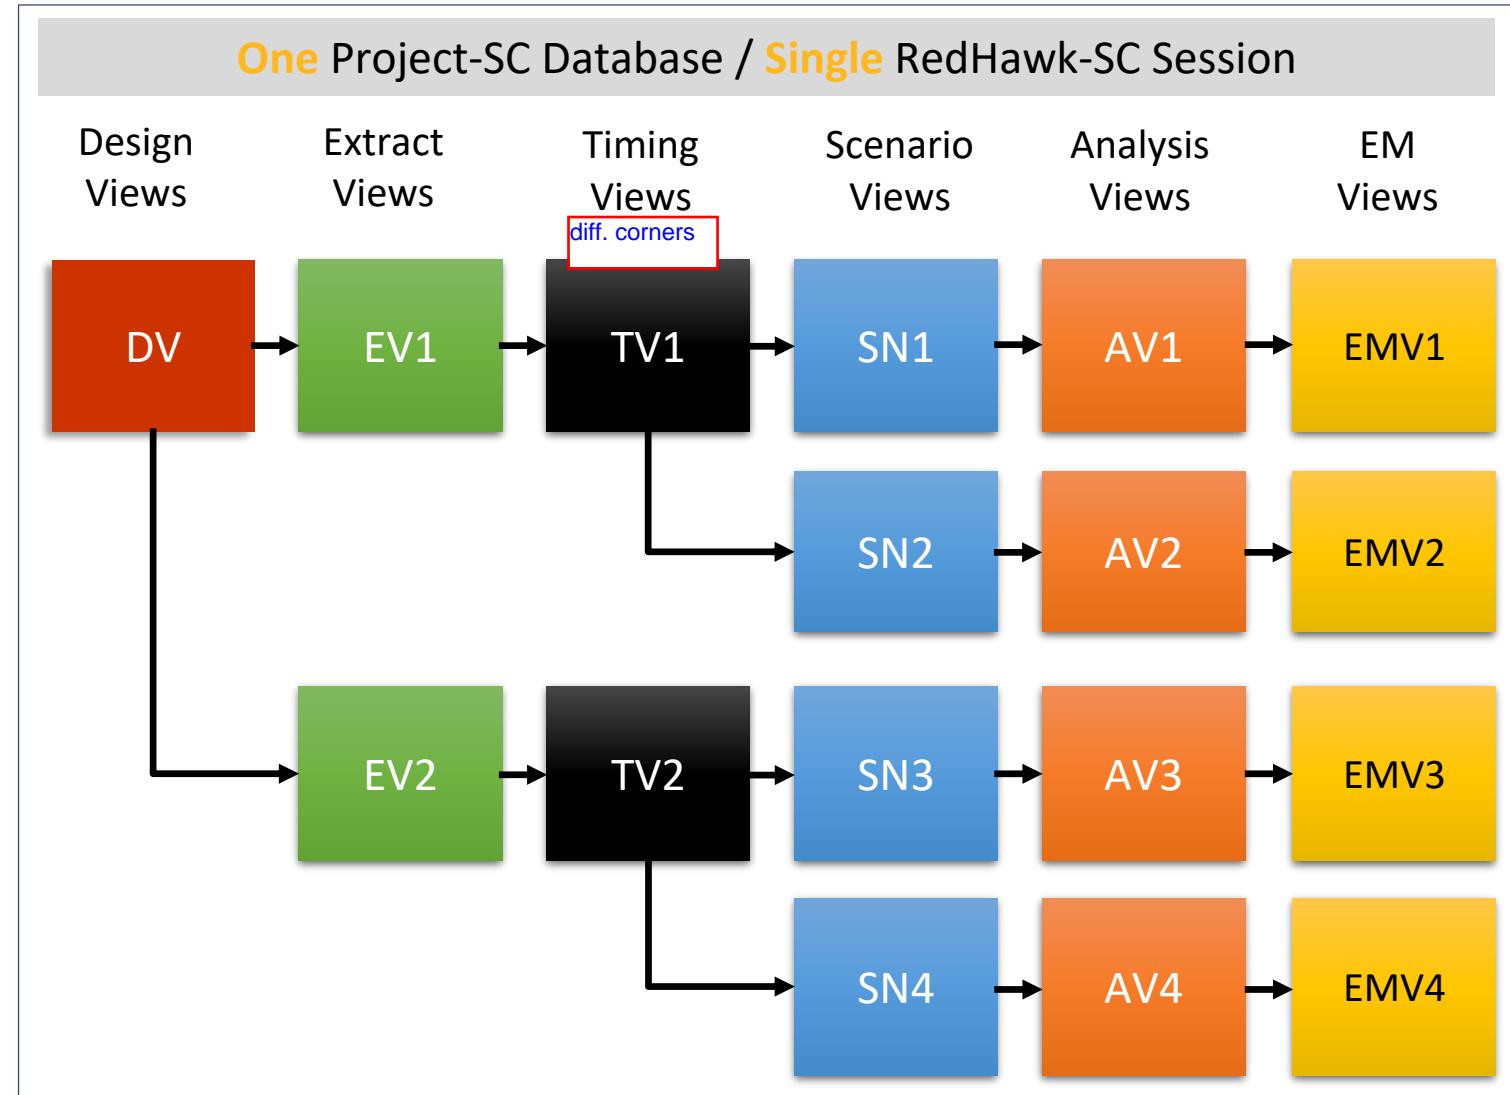


- Views are sub-directories inside SeaScape DB
- They are representations of a group of data
- Each view stores data from different stages of analysis
  - Write once, read many times
  - Automatic versioning is maintained when a view is over-written
  - Tag/directory name will have additional #v1, #v2... appended
  - Example tag='tv', on disk will be: ./<SeaScapeDB>/tv#v1)
- It is possible to restart the run from any view for faster TAT
- Views are python objects with query APIs to access their data

a view is written only Once, but can read many times. Once written, it is locked.

# Views - Example

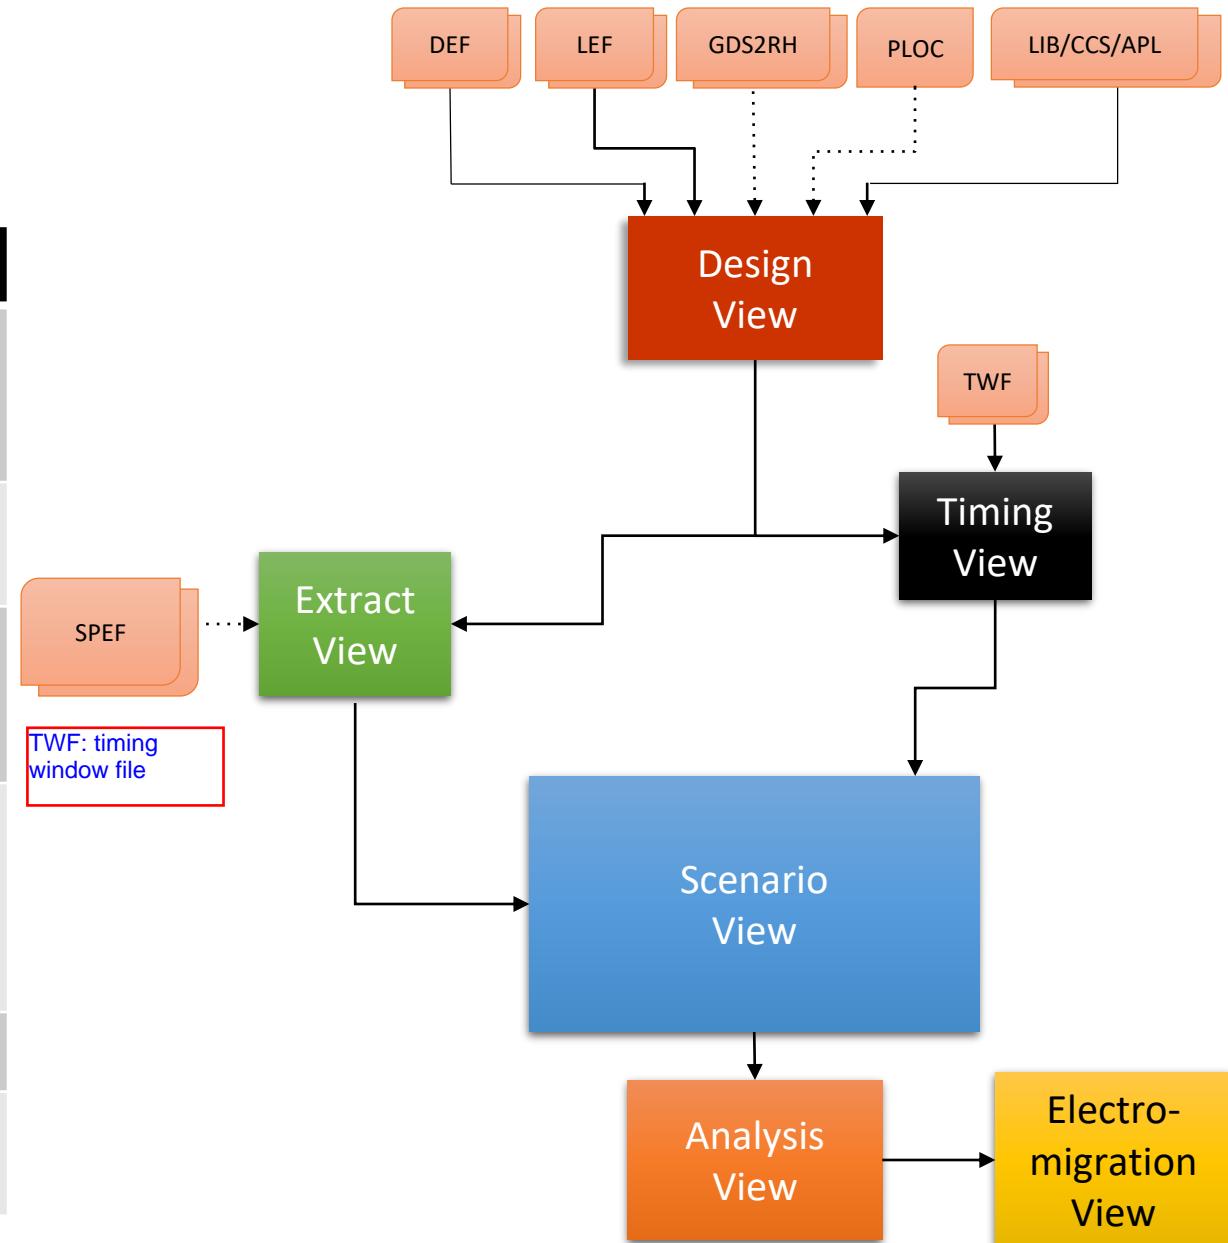
- **Single session can have views for:**
  - Different corners (Liberty and Extraction)
  - Different conditions (voltage, temperature)
  - Different modes of operation
  - Different vector sets
  - Different vector-less settings
  - Different conditions for IR drop versus EM
- **Example**
  - SN1, SN2 represent VCD & V-less scenarios
  - EV1 and EV2 represent 2 different extraction corners
- **Related views can be traced-back**
  - EMV4 ← AV4 ← SN4 ← TV2 ← EV2 ← DV



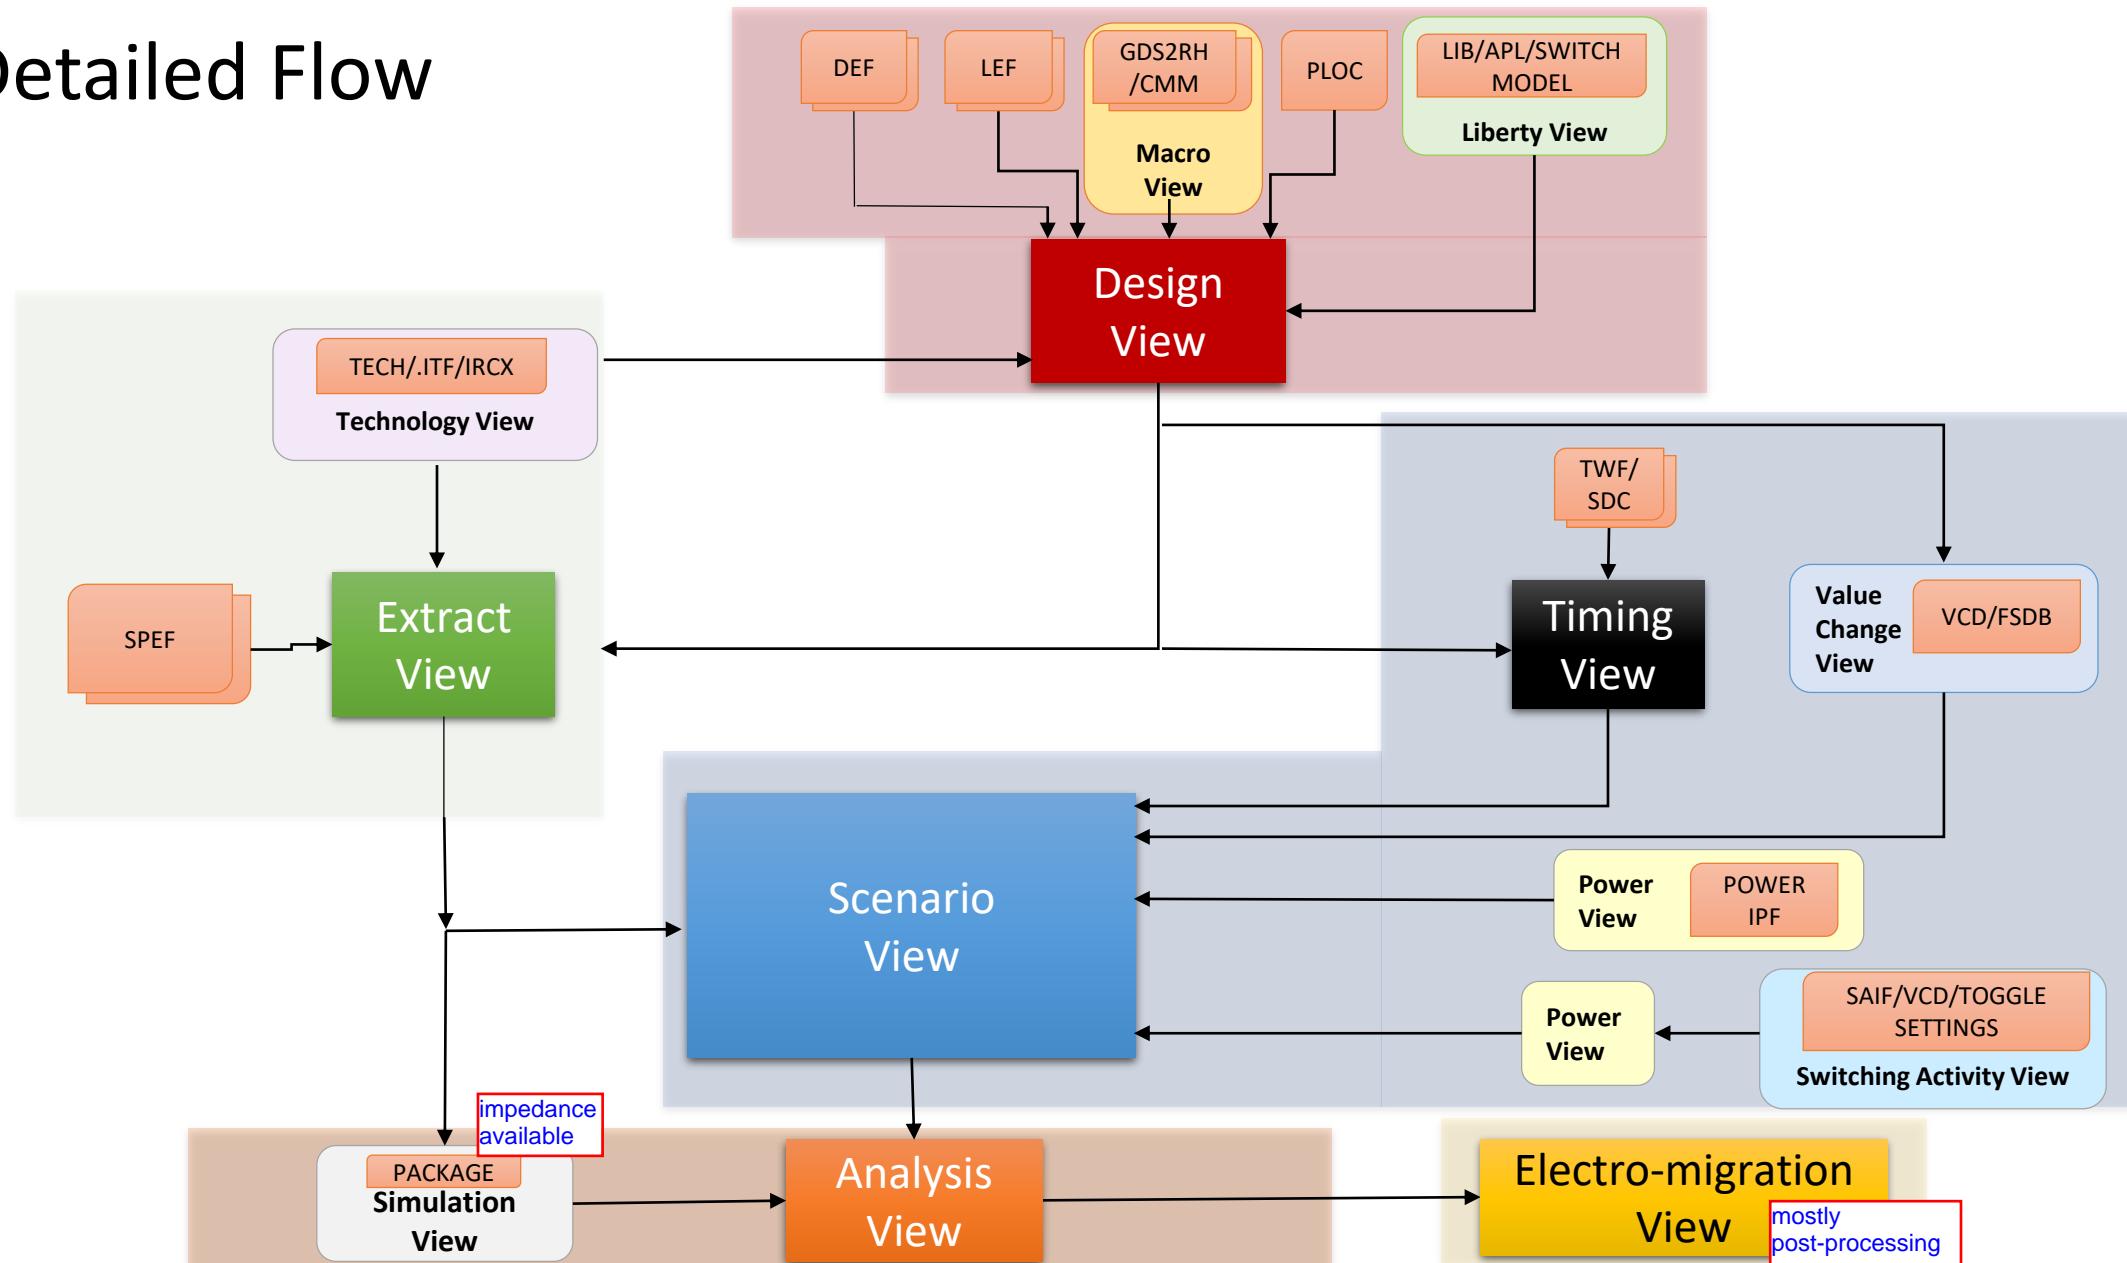
# Flow Overview

## Main SeaScape Views

View	Function
DesignView	<ul style="list-style-type: none"> <li>Stores netlist and layout information</li> <li>Library information (through LibertyView)</li> </ul>
ExtractView	<ul style="list-style-type: none"> <li>Stores RC extraction &amp; SPR data</li> <li>Imports and stores SPEF</li> </ul>
TimingView	<ul style="list-style-type: none"> <li>Stores imported TWF and SDC constraints</li> <li>Stores clock associations and slews.</li> <li>Computes timing graph.</li> </ul>
ScenarioView	<ul style="list-style-type: none"> <li>Calculates DC instance currents for <b>static scenario</b></li> <li>Stores logic events and transient instance currents for <b>dynamic scenario</b></li> </ul>
AnalysisView	<ul style="list-style-type: none"> <li>Stores result of simulation <b>IR sim.</b></li> </ul>
ElectromigrationView	<ul style="list-style-type: none"> <li>Reads in EM rules and produces EM results on every wire/via segment</li> </ul>



# More Detailed Flow



# Supplementary SeaScape Views

View	Function
<b>LibertyView</b>	<ul style="list-style-type: none"><li>Imports and stores Liberty, CCS models and APL models</li></ul>
<b>MacroView</b>	<ul style="list-style-type: none"><li>Read in macro models like gds2rh , Totem CMM models , etc</li></ul>
<b>TechView</b>	<ul style="list-style-type: none"><li>Stores extraction technology information from Ansys tech file or derived from ITF or iRCX</li></ul>
<b>ValueChangeView</b>	<ul style="list-style-type: none"><li>Stores imported FSDB/VCD.</li></ul>
<b>SwitchingActivityView</b>	<ul style="list-style-type: none"><li>Stores toggles – propagated/Unpropagated</li><li>Reads in toggle rate settings ,SAIF, VCD/FSDB etc</li></ul>
<b>PowerView</b>	<ul style="list-style-type: none"><li>Stores imported instance power file</li><li>Processes power and toggle scaling</li><li>Computes power based on toggles – propagated or unpropagated from SwitchingActivityView</li></ul>
<b>SimulationView</b>	<ul style="list-style-type: none"><li>Sets up for the analysis and reads in package models</li></ul>

# Launching RedHawk-SC

- **Batch mode execution example:**

- `<path_to_rhsc_installation>/bin/redhawk_sc <python_command_file>`

- **Interactive mode execution example:**

- `<path_to_rhsc_installation>/bin/redhawk_sc -i <python_command_file>`
  - It needs an `exit()` command in script or entered manually to exit the Python shell

- **Connecting to a live RedHawk-SC run:**

- RedHawk-SC allows querying of data/results from an active session, by remotely attaching to the session
  - Multiple users can attach to the same session from multiple machines for querying/viewing results
  - `<path_to_rhsc_installation>/bin/redhawk_sc -r <gp_dir>`

- **Execution Log Files:**

- All RHSC log files reside by default in `gp<>` folder
    - If the run is fired in the same directory, tool will create `gp.1`, `gp2` incrementally. '`latest(gp)`' link will point to the most recent `gp` directory
  - Main log file for RedHawk-SC would be `<gp_directory>/run.log` file.
  - Other important log files are `run_v.log` ( more detailed log file for master ), `Worker*.log` files for individual workers.

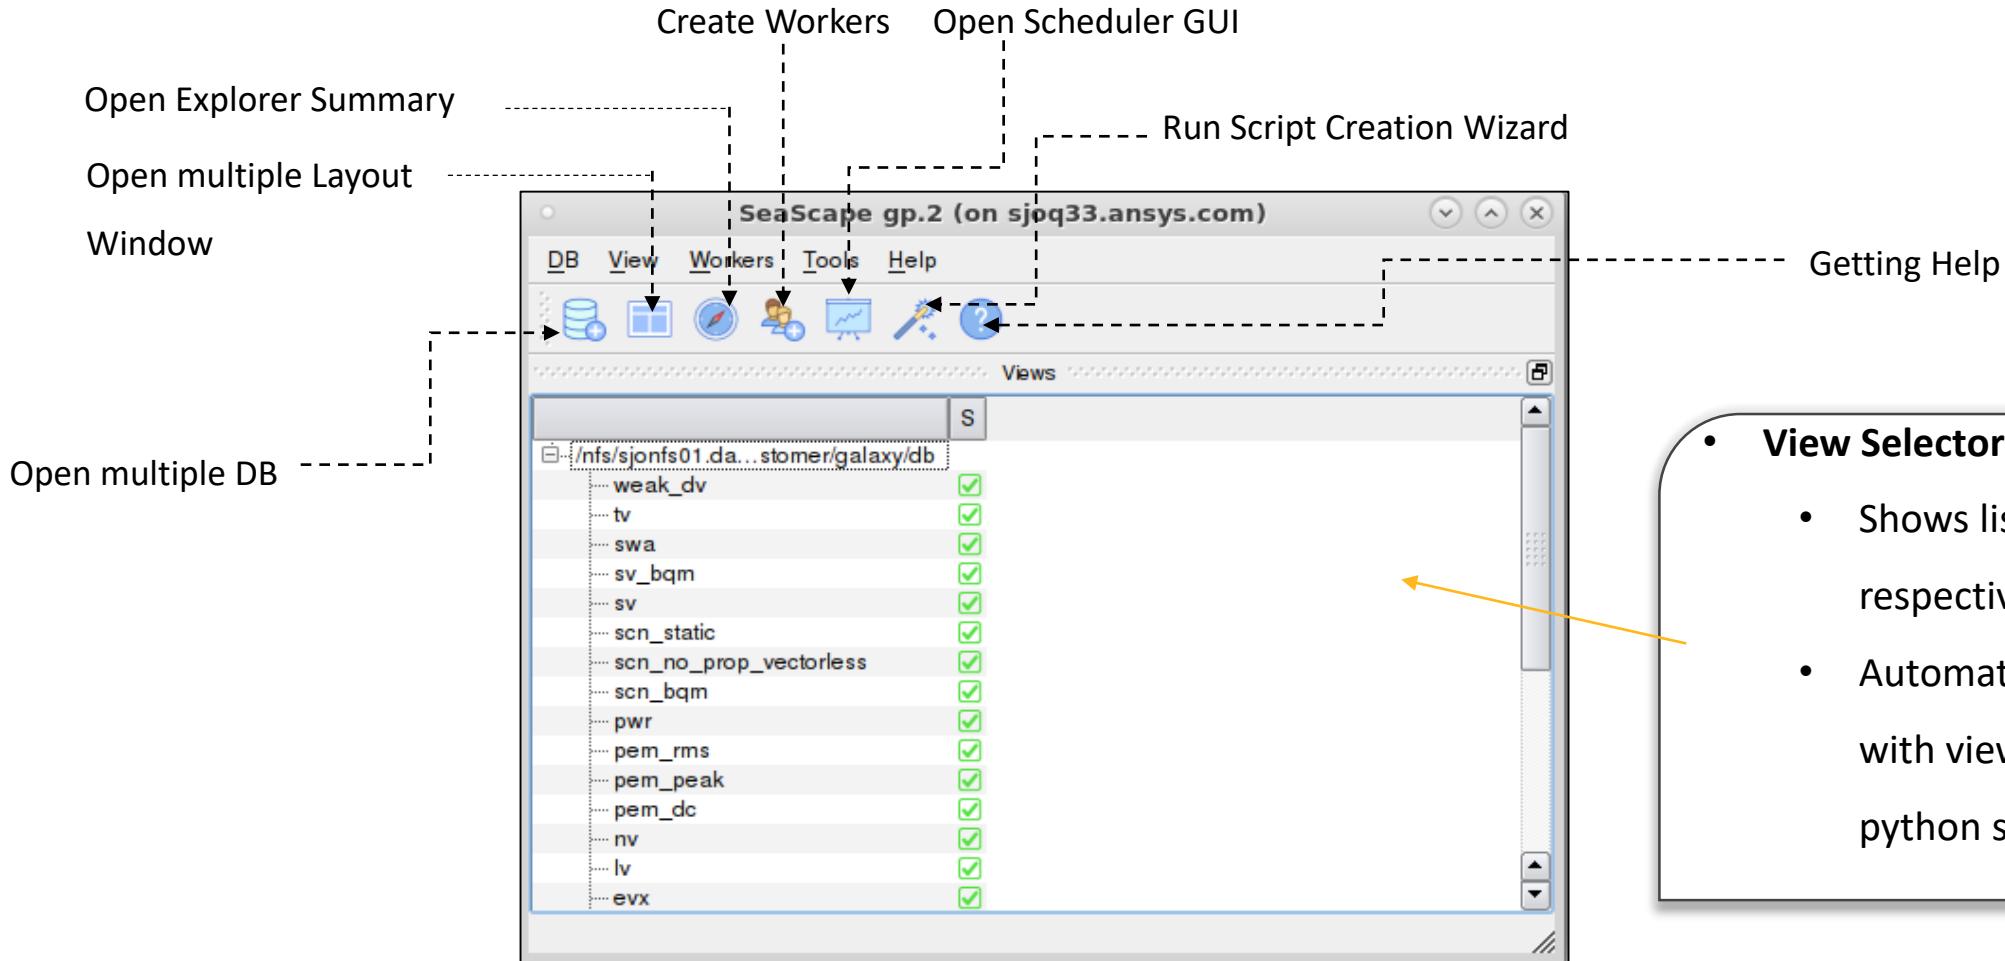
# SeaScape GUI Overview



# Launching RedHawk-SC Console

Run command : `<path_to_rhsc_installation>/bin/redhawk_sc <python_script> --console`

To invoke from RHSC shell , type : `open_console_window()`



- **View Selector**

- Shows list of open DBs and their respective views
- Automatically create variables with view names in the RH-SC python shell

# Launching Scheduler

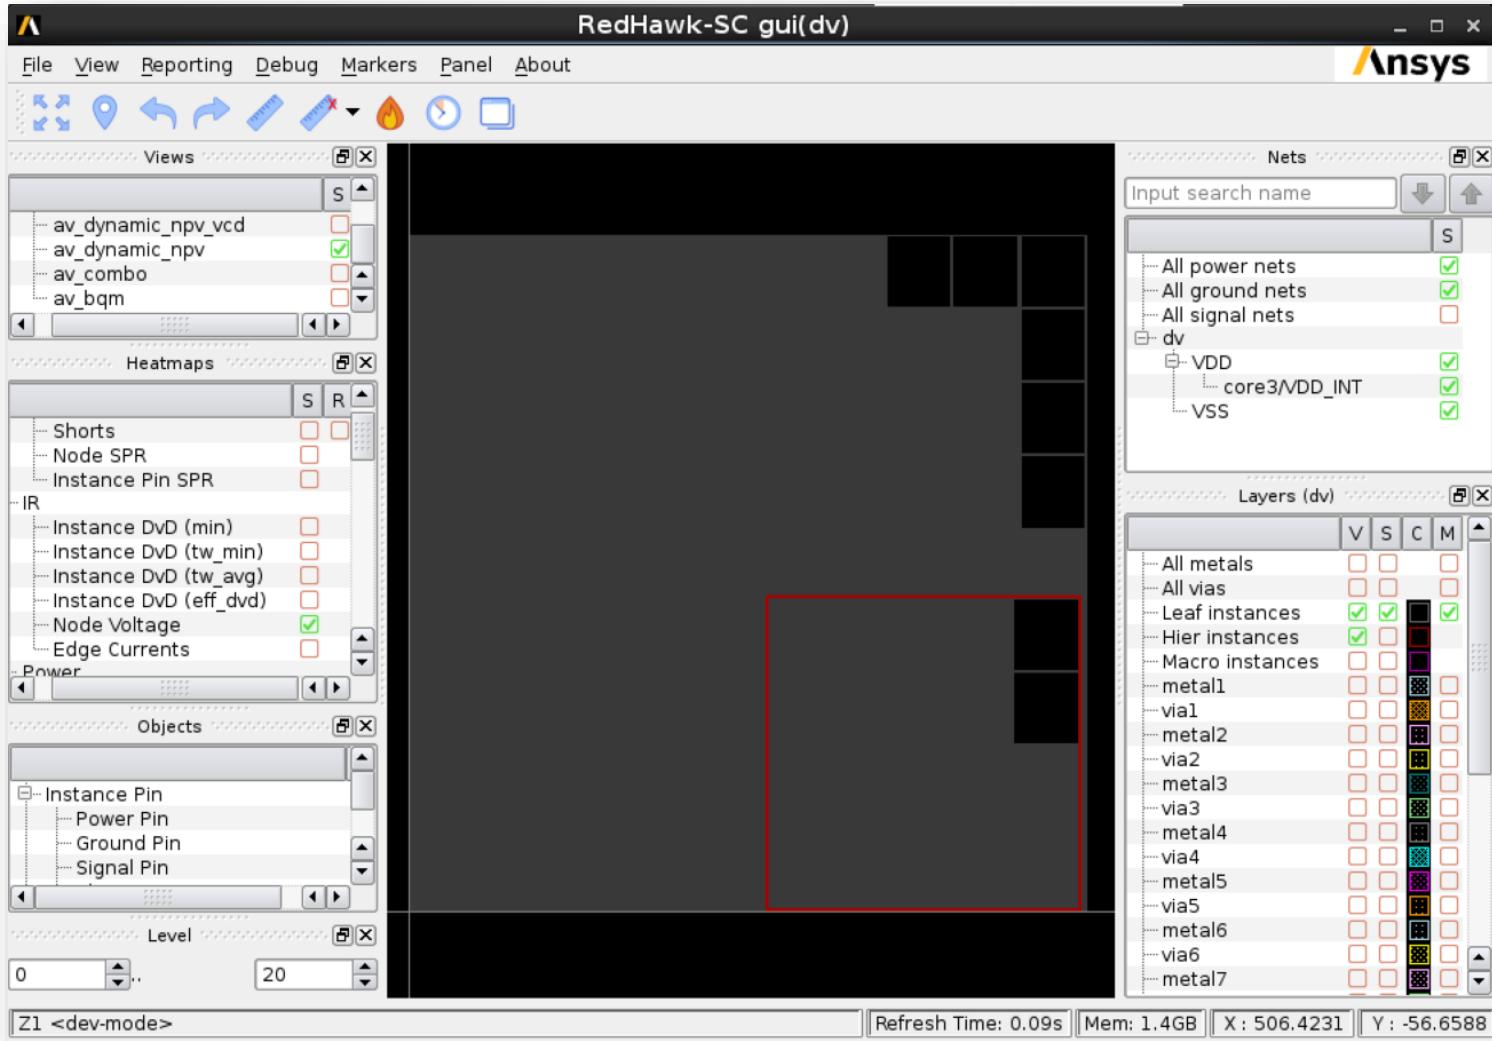
- Provides real time visualization of workers and their jobs
- Displays the jobs executed per worker (Y - axis) on the time axis (X - axis)
- Can be invoked from console or using the command '`open_scheduler window()`'
- For an ongoing or completed run , from the command line , we can invoke the scheduler gui by the command : `redhawk_sc -g -l -r <path_to_gp_directory>`



## FEATURES OF SCHEDULER GUI

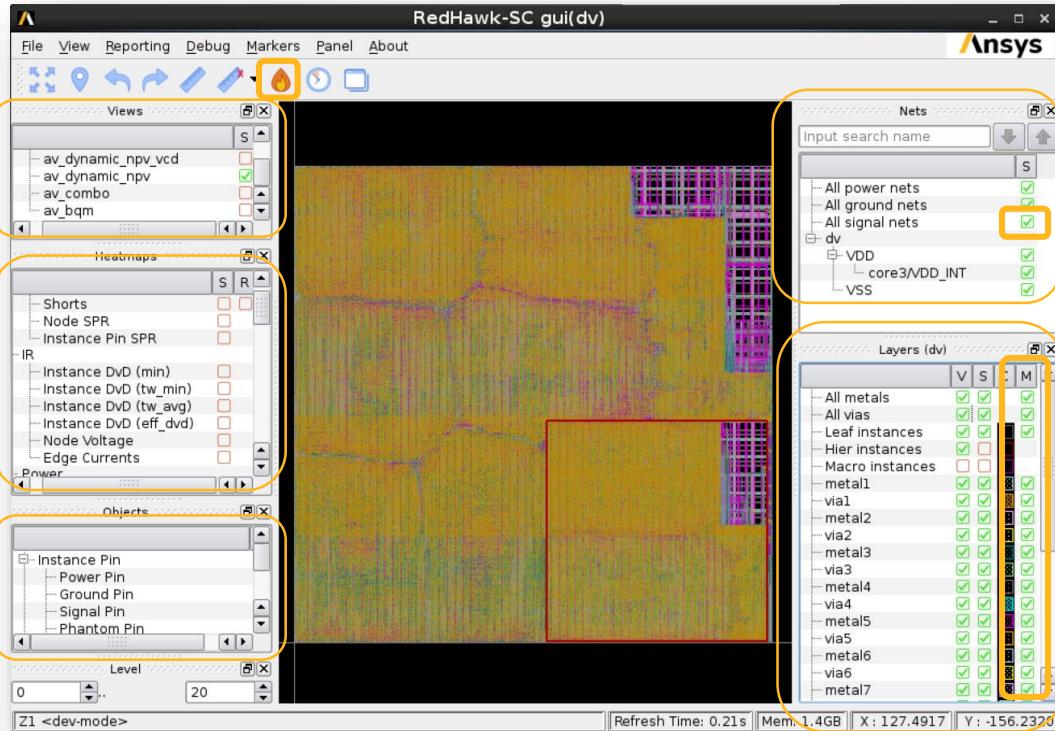
- **Worker info**
- **Host info**
- **Highlight Jobs using**
  - run-time
  - memory
- **Total number of workers, jobs**
  - name
- **Worker Stats Plots**
  - Worker memory
  - Machine load
  - Read ds rate
  - Total ds read
- **Identify network slowness/issues**
  - red boxes
  - using plots
- **Identify performance sticks \***

# RedHawk-SC™ GUI - Thin Client

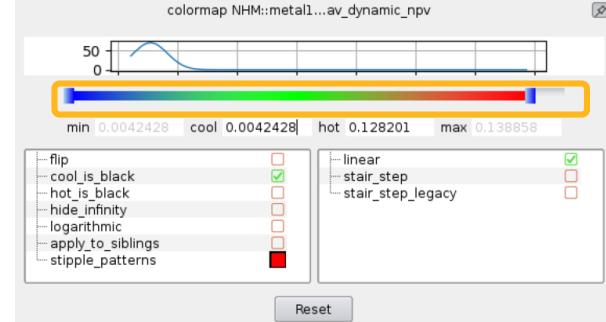
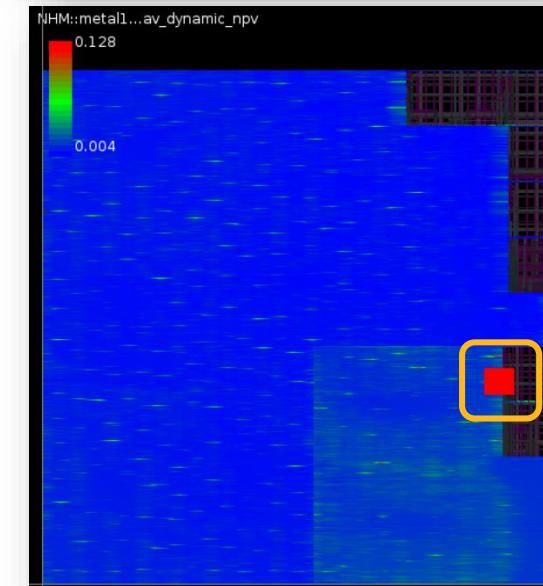
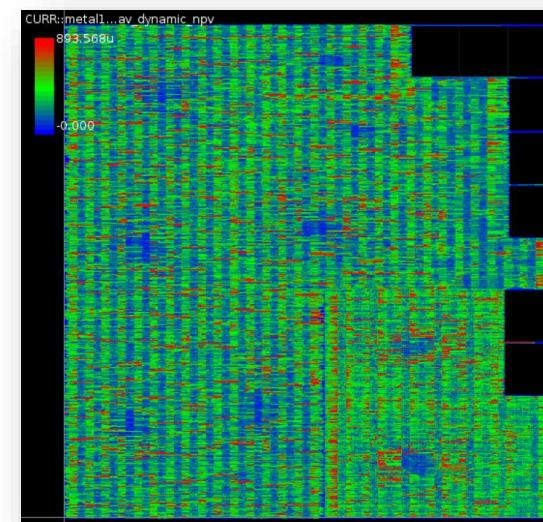


- GUI is a separate process that connects to master
- Multiple GUIs can be opened at same time with locked zoom
- It can run on a different host, or different network
- All data reads are on demand
- Very low memory consumption for viewing largest designs (~ 16 GB)

# RedHawk-SC™ GUI - Thin Client



- ✓ **VIEW SELECTOR** : viewing results from multiple views in same GUI
- ✓ **HEATMAP SELECTOR** : enabling multiple heatmaps together
- ✓ **OBJECT SELECTOR** for easy access to all PnR world objects – pins , vias , OBS etc
- ✓ **NET SELECTOR** : show PG and signal nets in single GUI
- ✓ **LAYER SELECTOR** : quick enabling/disabling layers
  - **MAP SELECTOR** ( M button ) for easy map access for instance based/layer based heatmaps



- ✓ **SLEEK SLIDER** to instantly change map thresholds

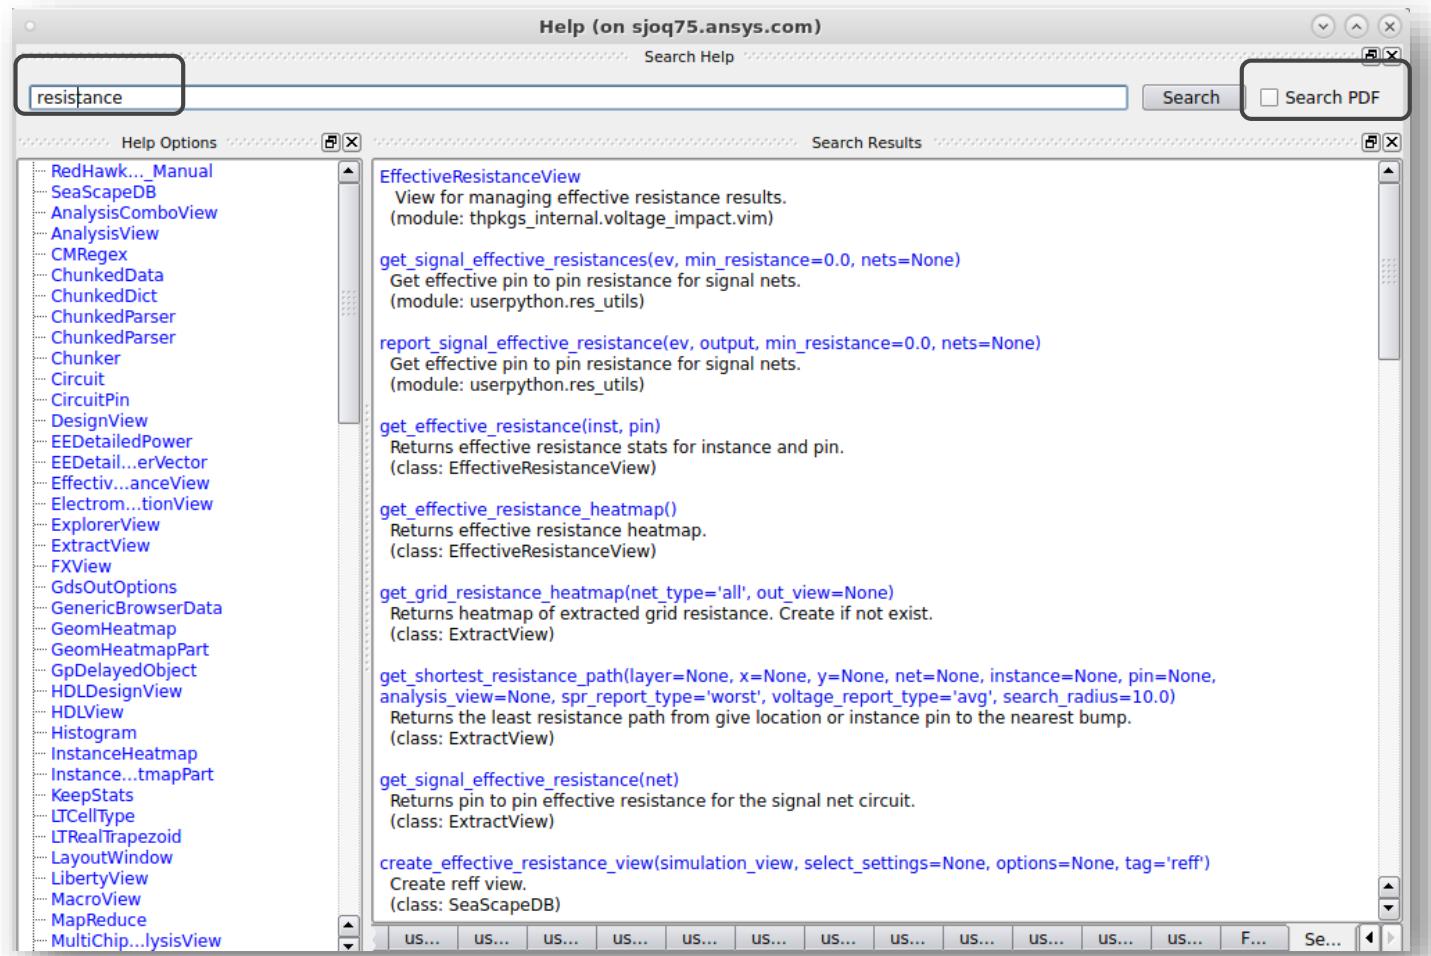
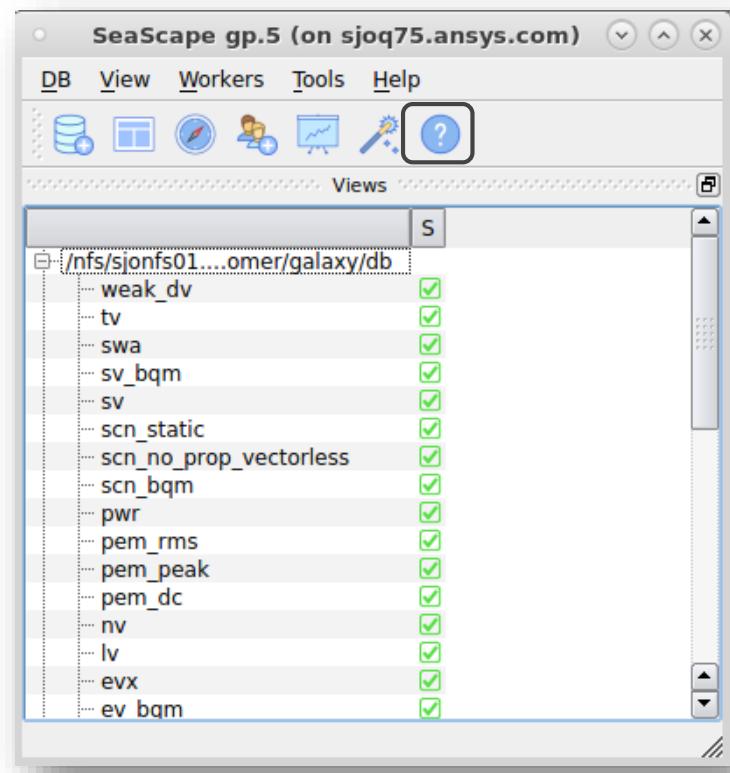


- ✓ **FLAME BUTTON** to quickly find hotspots above a threshold

# **Introduction to SC Help System & APIs**

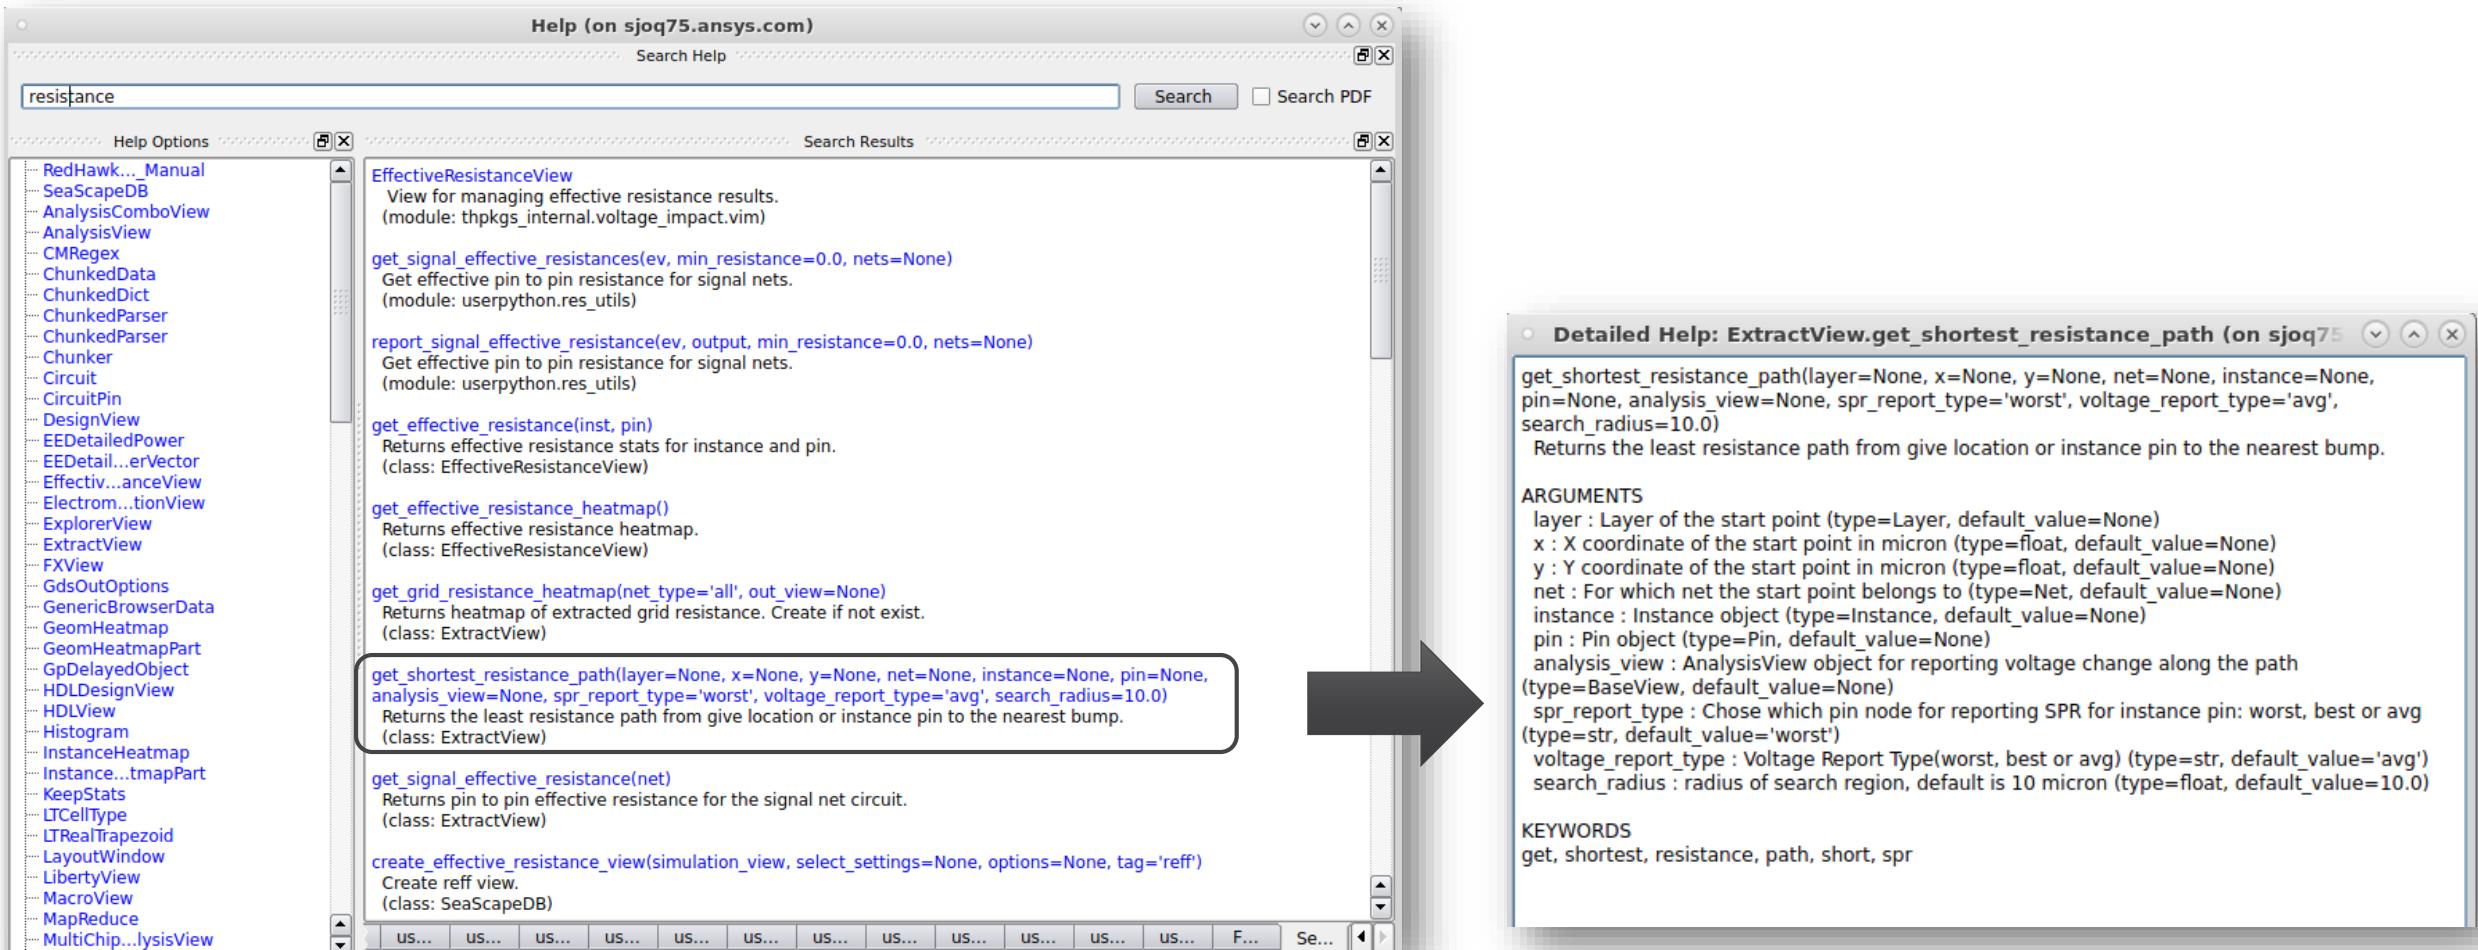


# Accessing Help From RedHawk-SC Console



- Open console and click on help button
  - Enter the item , for e.g resistance in the search box and get all results
  - Click on 'Search PDF' if you want to search inside RedHawk-SC documentation

# Accessing Help From RedHawk-SC Console



- Click on any search result to get more details

# Application Notes

```
indqa@sjoindqa04.ansys.com - /projs00/tomahawk/build/r20_R2/20.05.23_09.01.04/linux_x86_64_rhel6/Optimized/seascape_release/2020_R2.0.ENG/linux_x86_64_rhel6/doc/training  
- 96>  
ls AppNote_*  
AppNote_Built-in_Report.pdf  
AppNote_CPM.pdf  
AppNote_CTM_Creation_in_RedHawk-SC.pdf  
AppNote_ClockFX_RedHawk-SC.pdf  
AppNote_Controling_Macro_Switching_for_Dynamic_Analysis.pdf  
AppNote_Decap_Handling_in_RedHawk-SC.pdf  
AppNote_Diskspace_Management_in_Seascape.pdf  
AppNote_Early_Grid_Analysis_and_Robustness_Checks.pdf  
AppNote_FSDB_Replay_in_Fullchip_Dynamic_Simulation.pdf  
AppNote_In-Rush_Analysis.pdf  
AppNote_Introduction_to_Heatmaps.pdf  
AppNote_Macro_Handling_in_RedHawk-SC.pdf  
AppNote_Multi-Chip_Analysis_in_RedHawk-SC.pdf  
AppNote_NPV.pdf  
AppNote_Package_Board_Analysis.pdf  
AppNote_PowerProfileView.pdf  
AppNote_Presim_Handling_in_RedHawk-SC.pdf  
AppNote_RedHawk-SC_License_Management.pdf  
AppNote_Resiliency_in_RedHawk-SC.pdf  
AppNote_Rollup_Analysis.pdf  
AppNote_SEB.pdf  
AppNote_Scaling_Power_in_PowerView.pdf  
AppNote_Scheduler_GUI.pdf  
AppNote_SelfHeat_Analysis.pdf  
AppNote_Signal_EM.pdf  
AppNote_Time-step_Determination_in_RHSC.pdf  
AppNote_Using_Mapreduce_in_Redhawk-SC.pdf  
AppNote_VCD_Handling.pdf  
AppNote_Vectorless_Scan.pdf  
AppNote_Voltage_Stats_Report.pdf  
AppNote_Worker_Management_in_Redhawk-SC.pdf  
AppNote_ecos_in_RedHawk-SC.pdf
```



SeaScape gp.2 (on sjoindqa04.ansys.com)

Views

- RedHawk\_Manual
- SeaScapeDB
- AnalysisComboView
- AnalysisView
- CMBregex
- ChunkedData
- ChunkedDict
- ChunkedParser
- Chunker
- Circuit
- CircuitPin
- DesignView
- EEDetailedPower
- EEDetailed\_enVector
- EffectanceView
- Electrom...ionView
- ExplorerView
- ExtractView
- FXView
- GdsOutOptions
- GenericB...userData
- GeomHeatmap
- GeomHeatmapPart
- GpDelayedObject
- HDLDesignView
- HDLView
- Histogram
- InstanceHeatmap
- Instance\_tmapPart
- KeepStats
- TCellType
- TRRealTapezoid
- LayoutWindow
- LibertyView
- MCVoltageSigmas
- MacroView
- MapReduce

Help (on sjoindqa04.ansys.com)

Search Help

Search Results

Search PDF

RedHawk-SC\_User\_Manual

RedHawk-SC\_ReleaseNotes\_2020\_R1.0

AppNote\_CPM

AppNote\_Package\_Board\_Analysis

AppNote\_Presim\_Handling\_in\_RedHawk-SC

RedHawk-SC\_ReleaseNotes\_2020\_R1.2

RedHawk-SC\_ReleaseNotes\_2020\_R1.3

AppNote\_Package\_Board\_Analysis.pdf (on sjoindqa04.ansys.com)

File Edit View Go Help

Index

1. Introduction 1

2. Package and Board Models 1

3. Package Port Count Recommendation 11

4. Important functions and their usage 11

5. Known Restrictions 17

6. Conclusion 18

Ansys

DIE Analysis with Package and Board in RedHawk-SC

Version: 2020.02.05

Table of Contents

1. Introduction ..... 1

2. Package and Board Models ..... 1

    2.1 Simple Package RLC model ..... 1

    2.2 Package layout extracted Distributed RLCK spice netlist ..... 3

    2.3 S-parameter package and board model ..... 7

3. Package Port Count Recommendation ..... 11

4. Important functions and their usage ..... 11

5. Known Restrictions ..... 17

6. Conclusion ..... 18

1. Introduction

High quality models of off-chip RLC circuit elements such as the package and board are needed for

# Looking at Explorer Results

The screenshot shows the SeScape gp.5 interface on the left and the Explorer Summary window on the right. In the SeScape interface, the 'explorer' icon (a magnifying glass) is highlighted with a red box and an arrow pointing to the Explorer window.

**Explorer Summary (on sjoq36.ansys.com)**

**Design info**

via_instance_count	3.438M
power_nets	[Net('core3/VDD_INT'), Net('VDD')]
instance_count	1.165M
hierarchical_instance_count	2.0
ground_nets	[Net('VSS')]
geometry_count	3.276M

**Analysis info**

- Scenario Demand VSS (scn\_no\_prop\_vectorless)
- Scenario Demand VDD (scn\_no\_prop\_vectorless)
- Basic Demand Power (scn\_static)
- Ideal Voltage Levels (scn\_static)
- Clock Info (scn\_static)
- Simulation Parameters (av\_static)
- Voltage Results Summary (av\_static)
- Power per Net (av\_static)
- Instance Voltage Histogram
- Convergence (av\_static)
- Reduction Stats (av\_static)
- Factorization Stats (av\_static)
- Simulation Parameters (av\_dynamic)
- Voltage Results Summary (av\_dynamic)**
- Decap Stats (av\_dynamic)
- Power per Net (av\_dynamic)
- Convergence (av\_dynamic)
- Reduction Stats (av\_dynamic)
- Factorization Stats (av\_dynamic)
- Downloaded elements

	average	worst
instance compression	1.093	1.005
core3/VDD_INT drop/metal2	6.35m	12.205m
core3/VDD_INT drop/metal1	3.551m	94.825m
VSS drop/metal9	3.787m	9.472m
VSS drop/metal8	3.827m	10.945m
VSS drop/metal7	3.877m	12.1m
VSS drop/metal6	3.882m	12.379m
VSS drop/metal5	3.888m	12.84m
VSS drop/metal4	3.893m	13.537m
VSS drop/metal3	3.903m	14.001m

Click on explorer icon  
from console to open  
Explorer window

# Accessing Help From RedHawk-SC Python Shell

```
>>> help(DesignView)
CLASS
DesignView
    View for design data including layout and connectivity.

SEE ALSO
SeaScapeDB.create_design_view

FUNCTIONS
convert_to_id(obj)
    Converts name-based object or iterable to id-based object or list.

convert_to_name(obj)
    Converts id-based object or iterable to name-based object or list.

get_attributes(obj, include_pin_geoms=None, include_phantoms=None, include_clock_gate_pins=None,
include_di_checks_stats=None, only_data_types=None, include_data_types=None)
    Returns a dict of DesignView attributes for the specified object.

get_cell_data_coverage()
    Returns a dict of DesignView DI Summary.

... ... ...
```

# Understanding the API – SeaScapeDB Objects

- SeaScape organizes data by SeaScapeDB objects, the most common of which are:
  - **Layer, Cell, Instance, Pin, Net** , etc
  - Each of these objects has both a string and integer representation
  - In distributed processing and MapReduce, strings are highly inefficient
  - Conversion between string and integer can be done through :
    - `<DesignView>.convert_to_name(<object>)`
    - `<DesignView>.convert_to_id(<object>)`

```
>>> dv.convert_to_id(Instance('core0.regfile_data_memory.OAI22_X1_821'))  
>>> Instance(89172)  
>>> dv.convert_to_name(Instance(89172))  
>>> Instance('core0.regfile_data_memory.OAI22_X1_821')  
>>> dv.convert_to_id(Pin('VSS'))  
>>> Pin(5)
```

# Object Attribute Query Examples

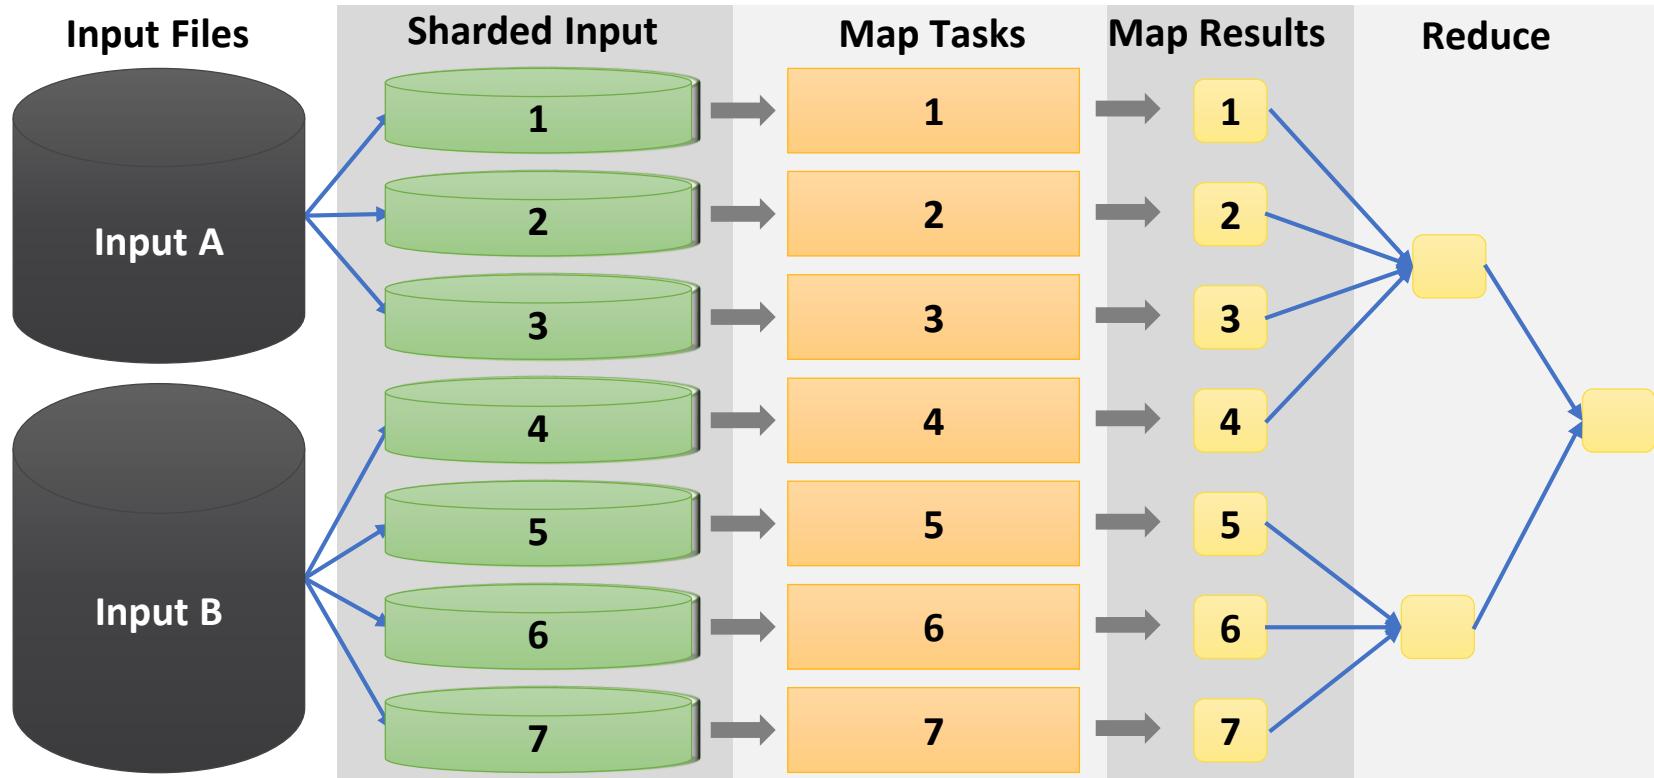
```
>>> dv.get_attributes(Instance("core1.regfile_program_memory.OAI21_X1_1777"))
>>> {'cell': Cell('OAI21_X2'),
     'cell_id': 135,
     'coord': RealCoord(145.92, 659.4),
     'instance_id': 202470,
     'is_leaf': True,
     'rotation': Rotation(RotationType.FS)}
```

```
>>> scn_no_prop_vectorless.get_attributes(Instance("core1.regfile_program_memory.OAI21_X1_1777"))
{'logic': {'clock_instance': False,
           'clock_tree_level': None,
           'clocks_reached': ['default_clock'],
           'frequency': 125000000.0,
           'load_cap': 2.8008853713451e-14,
           'logic_level': -1,
           ..... }}
```

# MapReduce Programming

# Introduction to MapReduce (MR)

**MapReduce (MR)** is a simplified model for processing and generating large data sets with a parallel, distributed algorithm on a cluster in an automated, easy to write, fashion.



**Disclaimer:** This section is going to cover only basics of MapReduce usage in RedHawk-SC.  
For detailed training, please refer to RedHawk-SC Modular Training on “Python APIs & MapReduce”

# Hands On MapReduce In SeaScape

- **What type of data am I going to analyze in MapReduce?**
  - Most common data which can be analyzed using MapReduce are sets of **Instances**
  - Also available for MapReduce are:
    - **Circuits** :: the actual extracted circuit
    - **Shapes** :: the geometries in the design
- **What type of data am I going to return?**
  - Most common data set returned from MapReduce is **list** and **dict**
  - Reduce operations are built-in within SeaScape
  - You can write your own custom reduce function also
  - Support for statistics, Waveforms, and others are available
- **What type of reduction will I need for my return?**
  - Any Python object that has '+' operator uses **reduce\_add**
  - For reducing dict, **reduce\_add\_dict** is available
  - Reduce\_add\_dict is the default reducer for MapReduce (when not specified)

# MapReduce Example: Find My Clock Instances In Design

```
def get_clk_inst_select(instances, scn):
    dv = scn.get_related_views(DesignView)[0]
    out = list()
    for inst in instances:
        is_clock = scn.get_attributes(inst)['logic'].get('clock_instance')
        if is_clock:
            out.append(inst)
    return out
```

## Serial Execution :

```
data = get_clk_inst_select(dv.get_instances('*'), scn=scn_no_prop_vectorless)
```

## Parallel Execution using MapReduce :

```
mm = MapReduce(dv)
mm.map_reduce(dv.get_mr_instances(), partial(get_clk_inst_select, scn=scn_no_prop_vectorless))
data = mm.get()
gp_print("Found ",len(data),' clock instances')
clki_selection = InstanceSelection(dv,data)
gui.add_layer(clki_selection,'clock_instance_map')
```

# Anatomy of a SeaScape MapReduce function

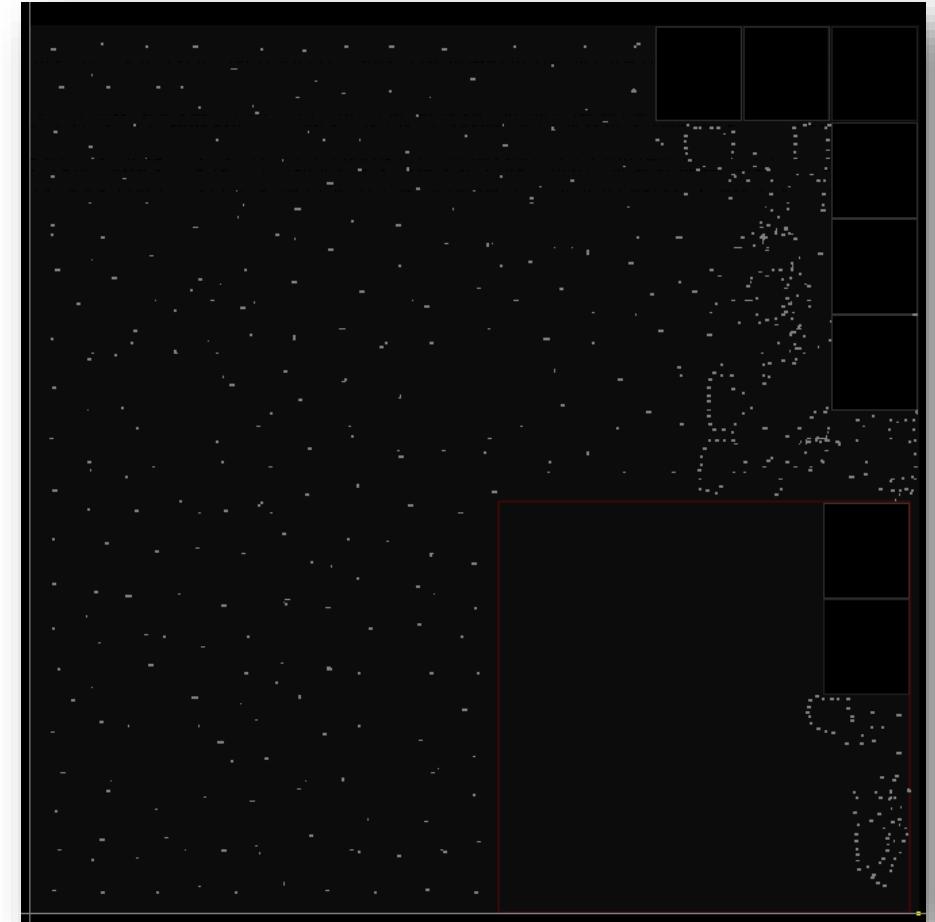
```
def map_func(<mapped data>, <arguments>):
    return <something>
    # map_func is the python job that will run in every worker
    # map_func will do some operations with a return data object (list, dict, etc.)
instances = dv.get_mr_instances()
# Get the sharded or mapped data , here gets sharded instances
# Sharded data can be Instances, Nets, Circuits etc
mm = MapReduce(dv)
# Informs the system what sharded data ( here Instances ) will be used in MapReduce.
# SeaScape internally keeps the instance data in sharded way for MapReduce operations
mm.map_reduce(<mapped data>,partial(<map_func>,<arguments by name>))
# Start the map and reduce operations .
data = mm.get()
# Do Something with the returned data from MapReduce
```

# MapReduce Output

```
>>> INFO<USER.000> Found 697 clock instances
```

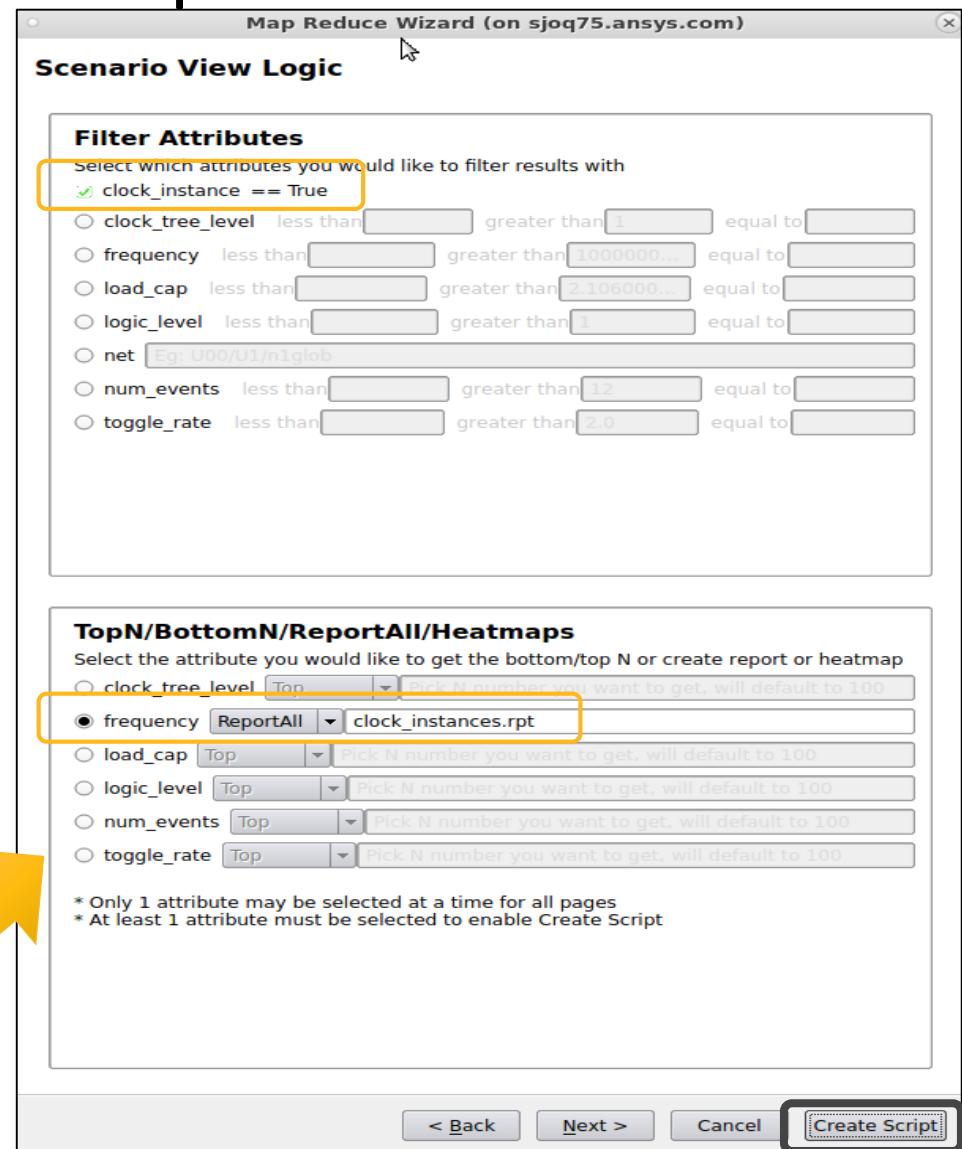
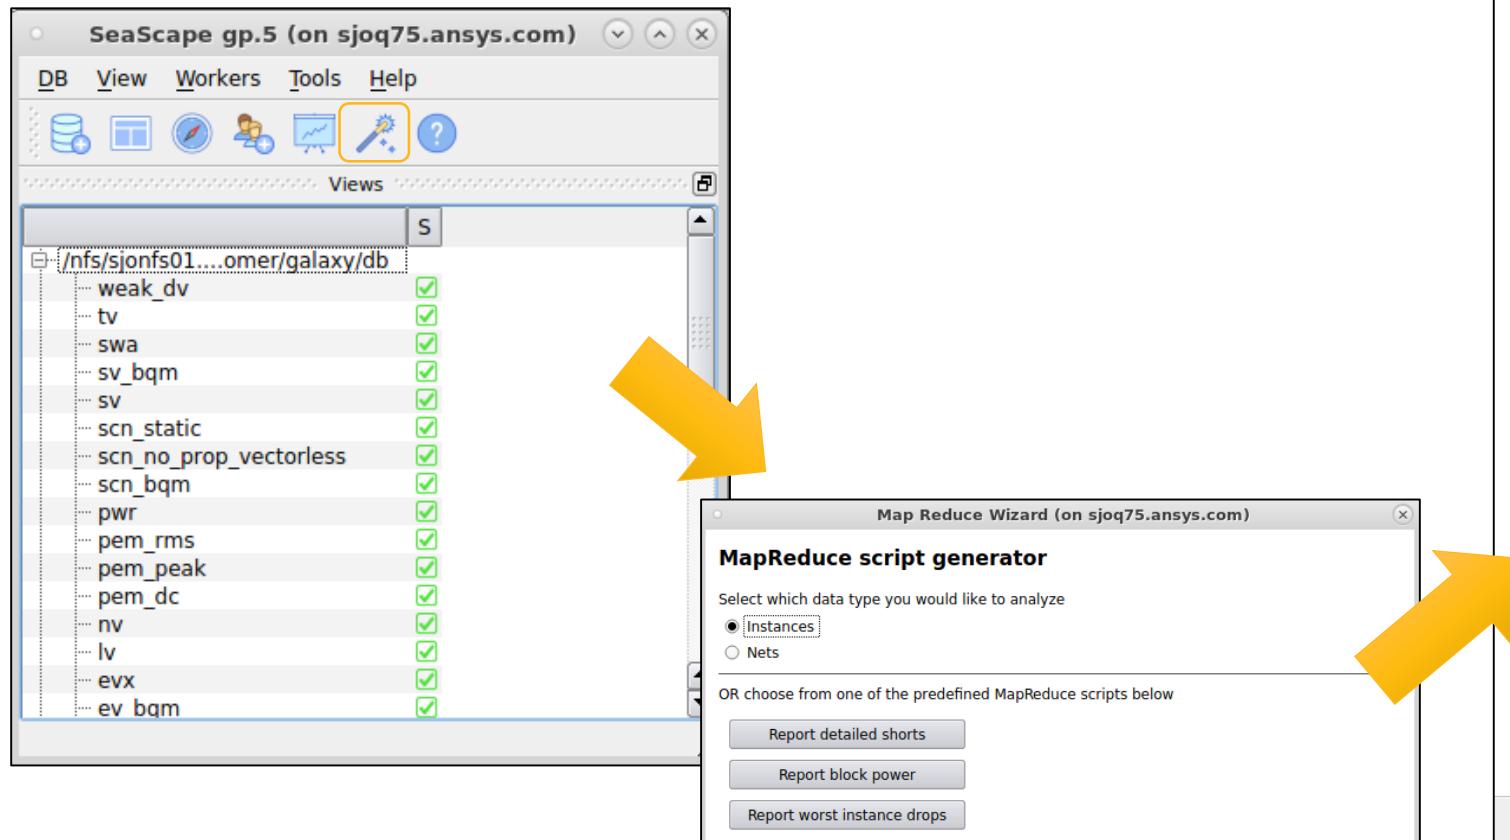
The parallel execution using MapReduce operations is much faster than serial execution on a large design

The clock instances from MapReduce are displayed in the GUI



# Using MapReduce Wizard To Run Python Scripts

- You don't need to be a Python expert to run MapReduce script in SC
- Use the MapReduce wizard to generate your Python script for most common use cases



# Using MapReduce Wizard To Run Python Scripts

Run script (on sjoq75.ansys.com)

File Click Here to run the script

```
from gp import *
open_scheduler_window(show_job_text=True)

def dv_filter(ii,dv_list):
    for dv_id, dv in enumerate(dv_list):
        dv_attributes = dv.get_attributes(ii)
        if dv_attributes.get('is_leaf') != True:
            return False
    return True

def scn_filter(ii,scn_list):
    for scn_id, scn in enumerate(scn_list):
        scn_attributes = scn.get_attributes(ii)
        if scn_attributes['logic'].get('clock_instance') != True:
            return False
    return True

def get_ii_frequency(ii,scn_list):
    stats = StdStatsDouble()
    for scn in scn_list:
        scn_attributes = scn.get_attributes(ii)
        frequency = scn_attributes['logic'].get('frequency')
        if frequency is not None:
            stats.add_sample(frequency)
    if stats.get_count() > 0:
        return (ii,stats.get('avg'))
    else:
        return None
```

Python script is automatically generated by MR Wizard



File Edit Tools Syntax Buffers Window Help

clock\_instances.rpt (/nfs/sjona/

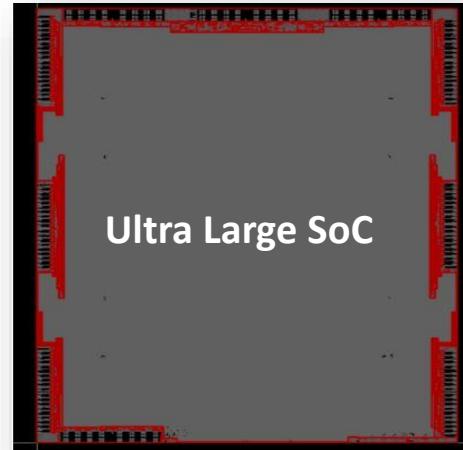
```
cts_inv_583961390 125000000.0
cts_inv_584361394 125000000.0
cts_inv_584461395 125000000.0
cts_inv_584961400 125000000.0
cts_inv_585261403 125000000.0
cts_inv_585361404 125000000.0
cts_inv_585461405 125000000.0
cts_inv_586361414 125000000.0
cts_inv_587161422 125000000.0
cts_inv_587261423 125000000.0
cts_inv_587761428 125000000.0
cts_inv_587861429 125000000.0
cts_inv_589861449 125000000.0
cts_inv_600161552 125000000.0
cts_inv_600461555 125000000.0
```

# Using SeaScape Platform



# Performance of Tools Using SeaScape Platform

N7            775 mm<sup>2</sup>            54B            66B  
TSMC        Design            Transistors        Nodes



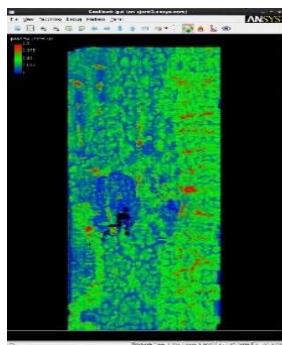
Analysis Type	Static + Dynamic
Total Run-time	41hrs
#Workers	2400
Peak worker memory	22GB

*RedHawk-SC delivers reliable big data compute over 100,000+ core-hours*

# An Example Of Big Data Analytics Application Using SeaScape Platform

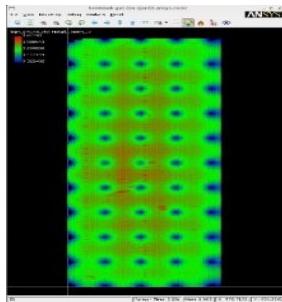
RAIL DATA

CURRENT CROWDING



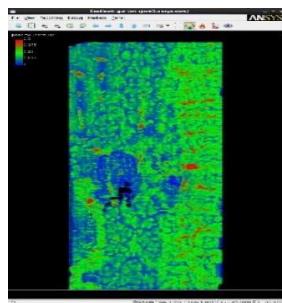
RAIL DATA

RESISTANCE



TIMING DATA

SLACKS



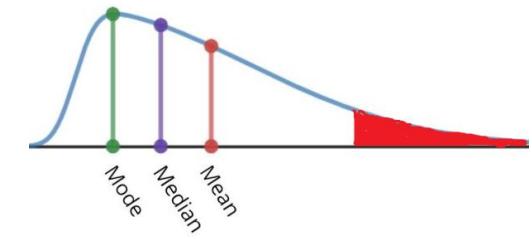
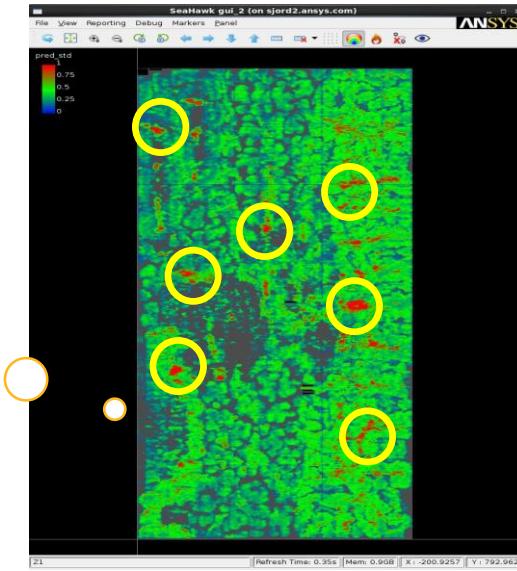
PHYSICAL  
DESIGN DATA  
PER LAYER  
CONGESTION

LIBRARY DATA  
CELL VOLTAGE  
SENSITIVITY

## Example Analytics

highlight instances which are timing critical, with high PG resistance and placed in high current crowding area

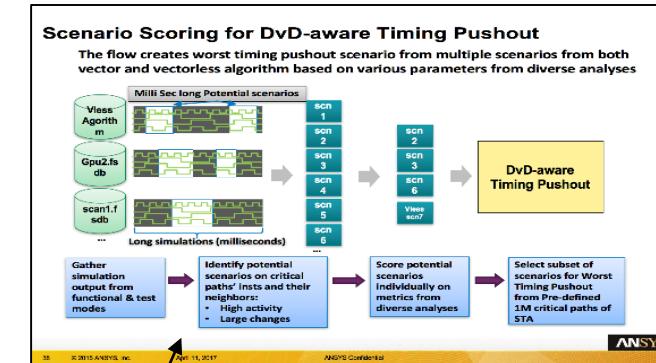
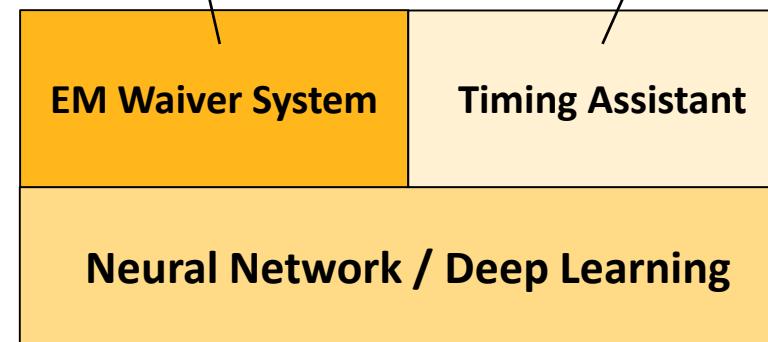
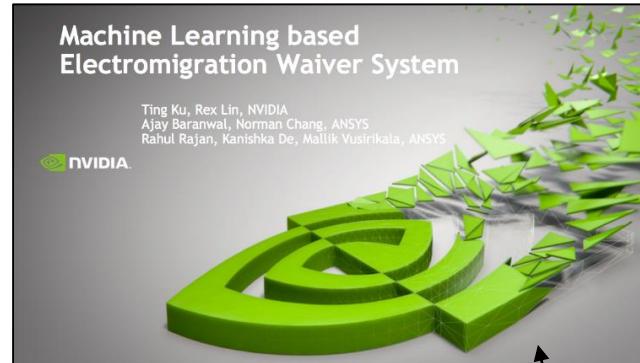
Multi-metric  
Analytics



Statistical distribution of outliers from SC data-analytics can help in finding design weakness

# SeaScape Platform

Integrated ML Stack for Scientific Computing



## SeaScape Big Data Platform

Actionable Analytics

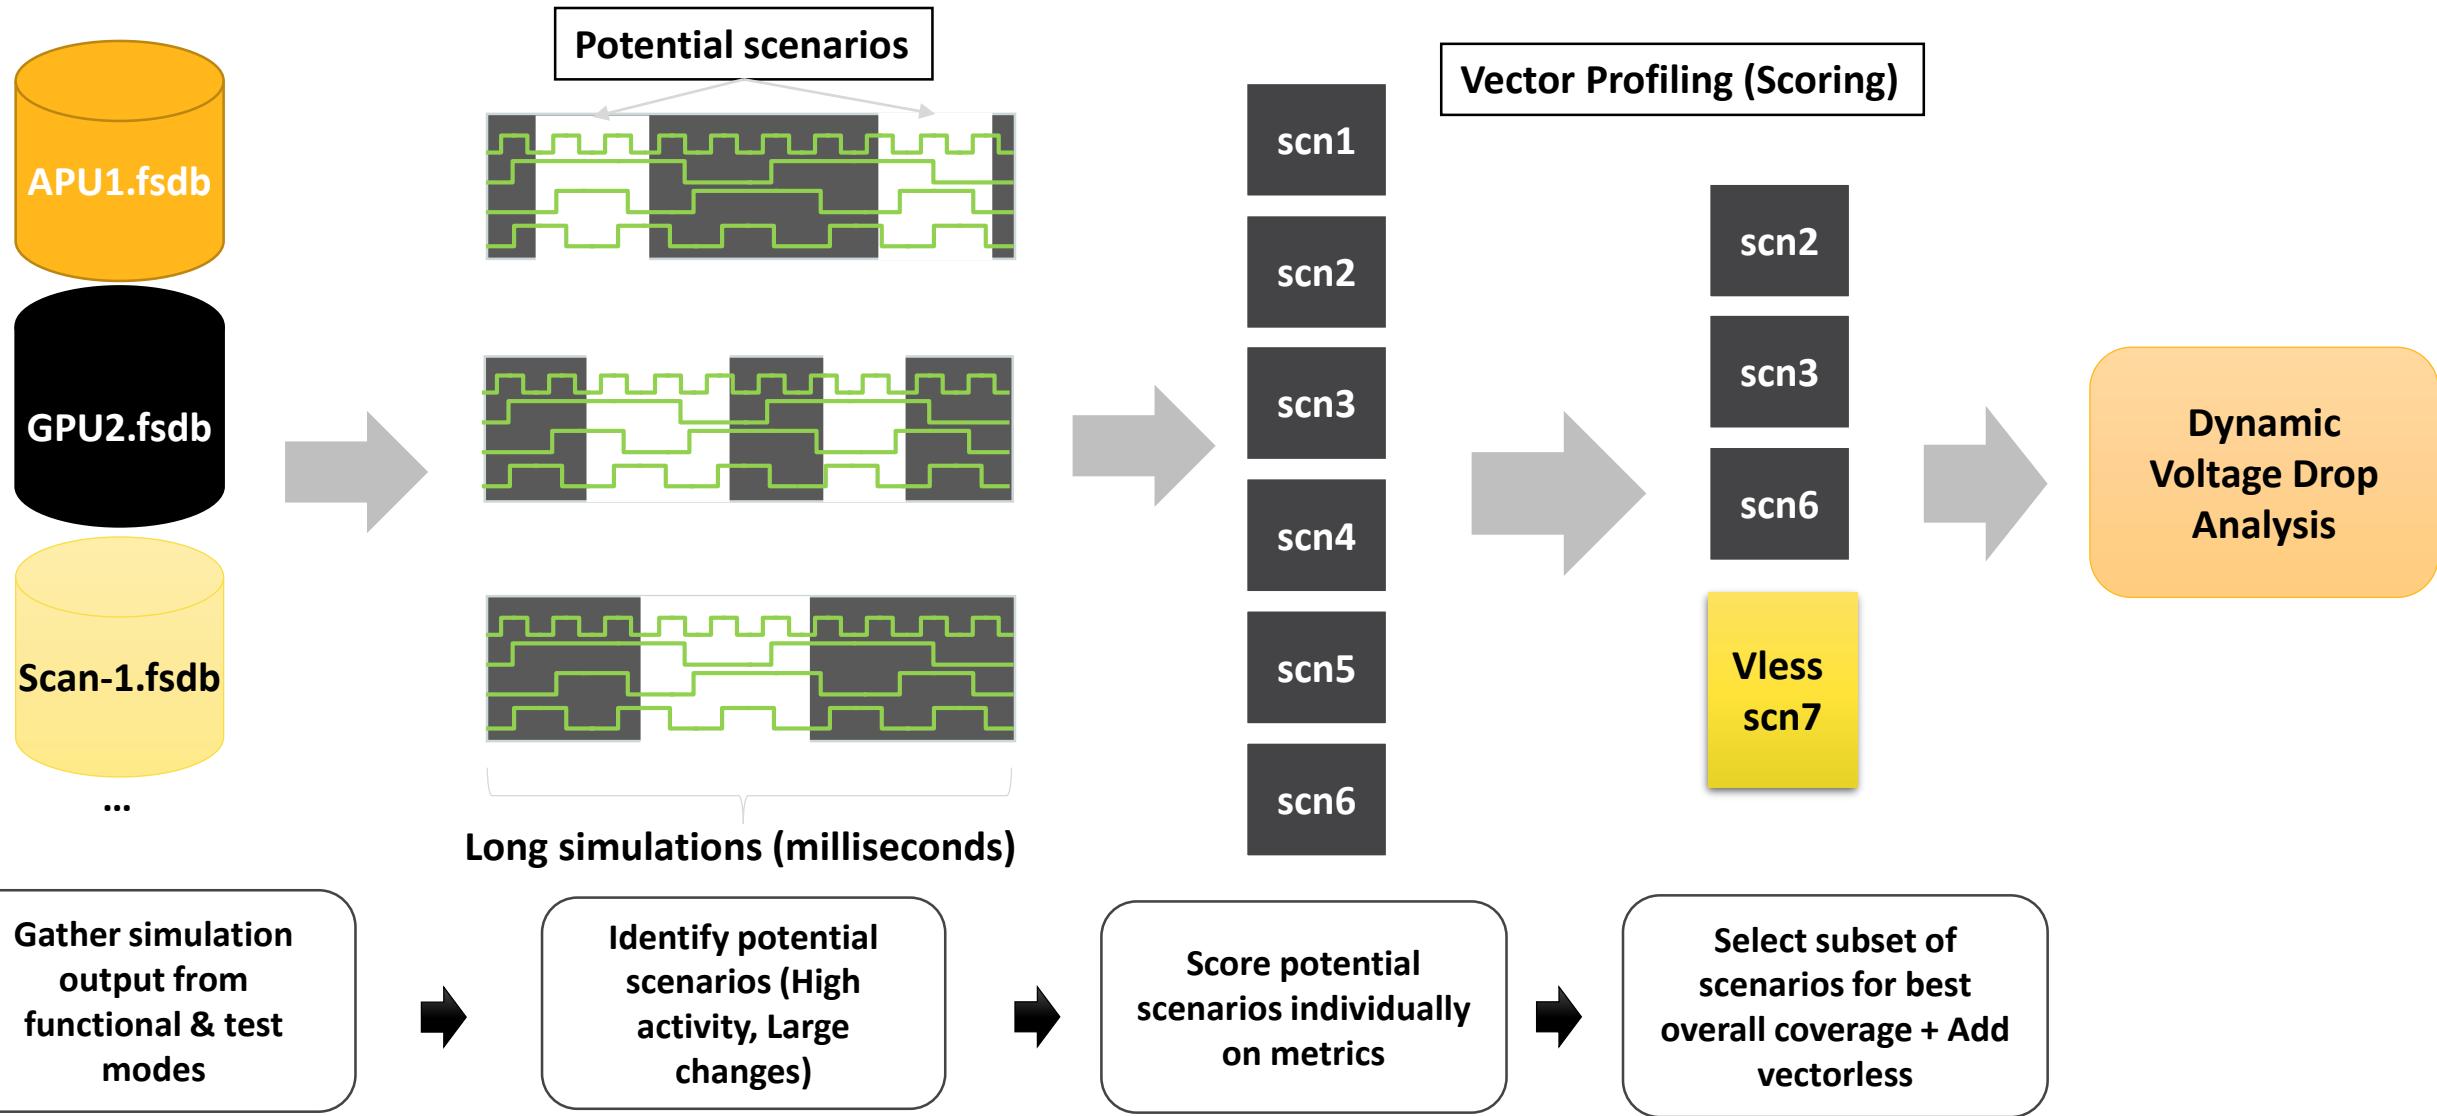
Machine Learning / AI\*

Elastic Compute

Visualization & Debug

\* Google TensorFlow Numpy K-Means Clustering, ...

# Scenario Scoring – Mix Of Vector + Vectorless Approaches



The Ansys logo consists of the word "Ansys" in a bold, black, sans-serif font. To the left of the "A", there is a graphic element: a thick yellow diagonal bar above a thick black diagonal bar, which together form a stylized 'A' shape.

Ansys

